

## ✗ FraudShield Assistant – Final Project (ITAI 2377)

### 1. Introduction

### 2. Data Collection & Preprocessing

### 3. Feature Engineering

### 4. Model Training & Selection

### 5. FraudShield Assistant Implementation

### 6. Evaluation

### 7. Reflection & Future Work

### 8. Team Contributions

#### Imports & Settings

```
1 from sklearn.metrics import (
2     classification_report,
3     confusion_matrix,
4     roc_auc_score,
5     average_precision_score,
6 )
7
8 import numpy as np
9 import pandas as pd
10
11 import matplotlib.pyplot as plt
12 import seaborn as sns
13
14 from sklearn.model_selection import train_test_split
15 from sklearn.preprocessing import StandardScaler, OneHotEncoder
16 from sklearn.compose import ColumnTransformer
17 from sklearn.pipeline import Pipeline
18 from sklearn.metrics import (
19     classification_report,
20     confusion_matrix,
21     roc_auc_score,
22     average_precision_score,
23     precision_recall_curve,
24 )
```

```
25
26 from imblearn.pipeline import Pipeline as ImbPipeline
27 from imblearn.over_sampling import SMOTE
28
29 from sklearn.linear_model import LogisticRegression
30 from sklearn.ensemble import RandomForestClassifier
31 from sklearn.ensemble import IsolationForest
32 from xgboost import XGBClassifier
33
34 import joblib
35
36 RANDOM_SEED = 42
37 np.random.seed(RANDOM_SEED)
38 plt.rcParams["figure.figsize"] = (8, 5)
```

## 1. Introduction

# 1. Introduction

In this project, we implement **FraudShield**, a domain-specific AI assistant that supports credit card fraud detection.

In our midterm project, we:

- Explored a synthetic credit card fraud dataset.
- Engineered domain-relevant features such as customer age, transaction time patterns, distance between customer and merchant, and spending velocity.
- Trained and evaluated several models (Logistic Regression, Random Forest, XGBoost, and Isolation Forest) under severe class imbalance.

In this final project, we:

- Organize the full data preprocessing and feature engineering pipeline.
- Perform model training and selection.
- Wrap the best-performing model in a simple **assistant interface** that:
  - scores transactions,
  - flags suspicious ones,
  - and provides a brief explanation.

We also evaluate the assistant using appropriate metrics for imbalanced classification and reflect on challenges and future work.

## ▼ 2. Data Collection & Preprocessing

## 2.1 Dataset & Domain

We use a synthetic credit card fraud dataset (`fraudTest_*.csv`) containing transaction records with:

- **Transaction details:** `trans_date_trans_time`, `category`, `merchant`, `amt`, etc.
- **Customer details:** `gender`, `city`, `state`, `dob`, geographic coordinates (`lat`, `long`).
- **Merchant location:** `merch_lat`, `merch_long`.
- **Target label:** `is_fraud` (1 = fraudulent, 0 = legitimate).

The dataset is **highly imbalanced**: only a small fraction of transactions are fraudulent.

```
1 # Loading Data
2 from google.colab import files
3
4 # Upload the dataset file (e.g., fraudTest_1000.csv)
5 uploaded = files.upload()
6
7 # Replace with your actual filename if different
8 DATA_PATH = list(uploaded.keys())[0]
9 print("Using data file:", DATA_PATH)
10
11 df = pd.read_csv(DATA_PATH)
12 df.head()
```

Choose Files No file chosen Upload widget is only available when the cell has been executed in the current browser session. Please rerun this cell to enable.

Saving fraudTest.csv to fraudTest (1).csv  
Using data file: fraudTest (1).csv

	Unnamed: 0	trans_date_trans_time	cc_num	merchant	category
0	0	2020-06-21 12:14:25	2291163933867244	fraud_Kirlin and Sons	personal_c...
1	1	2020-06-21 12:14:33	3573030041201292	fraud_Sporer-Keebler	personal_c...
2	2	2020-06-21 12:14:53	3598215285024754	fraud_Swaniawski, Nitzsche and Welch	health_fit...
3	3	2020-06-21 12:15:15	3591919803438423	fraud_Haley Group	misc_c...
4	4	2020-06-21 12:15:17	3526826139003047	fraud_Johnston-Casper	tr...

5 rows × 23 columns

```
1 # Basic Info & Nulls
2 df.info()
3 df.isnull().sum().sort_values(ascending=False).head(20)
```



```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 555719 entries, 0 to 555718
Data columns (total 23 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   Unnamed: 0        555719 non-null   int64  
 1   trans_date_trans_time 555719 non-null   object  
 2   cc_num            555719 non-null   int64  
 3   merchant          555719 non-null   object  
 4   category          555719 non-null   object  
 5   amt               555719 non-null   float64 
 6   first              555719 non-null   object  
 7   last              555719 non-null   object  
 8   gender             555719 non-null   object  
 9   street             555719 non-null   object  
 10  city               555719 non-null   object  
 11  state              555719 non-null   object  

```

```
Class Distribution
1 target_col = "is_fraud"
2 print(df[target_col].value_counts())
3 print("\nFraud rate:", df[target_col].mean())
4 sns.countplot(x=df[target_col])
5 plt.title("Class Distribution: is_fraud")
6 plt.show()

19  unix_time          555719 non-null   int64  
20  merch_lat          555719 non-null   float64 
21  merch_long         555719 non-null   float64 
22  is_fraud           555719 non-null   int64  
dtypes: float64(5), int64(6), object(12)
memory usage: 97.5+ MB
```

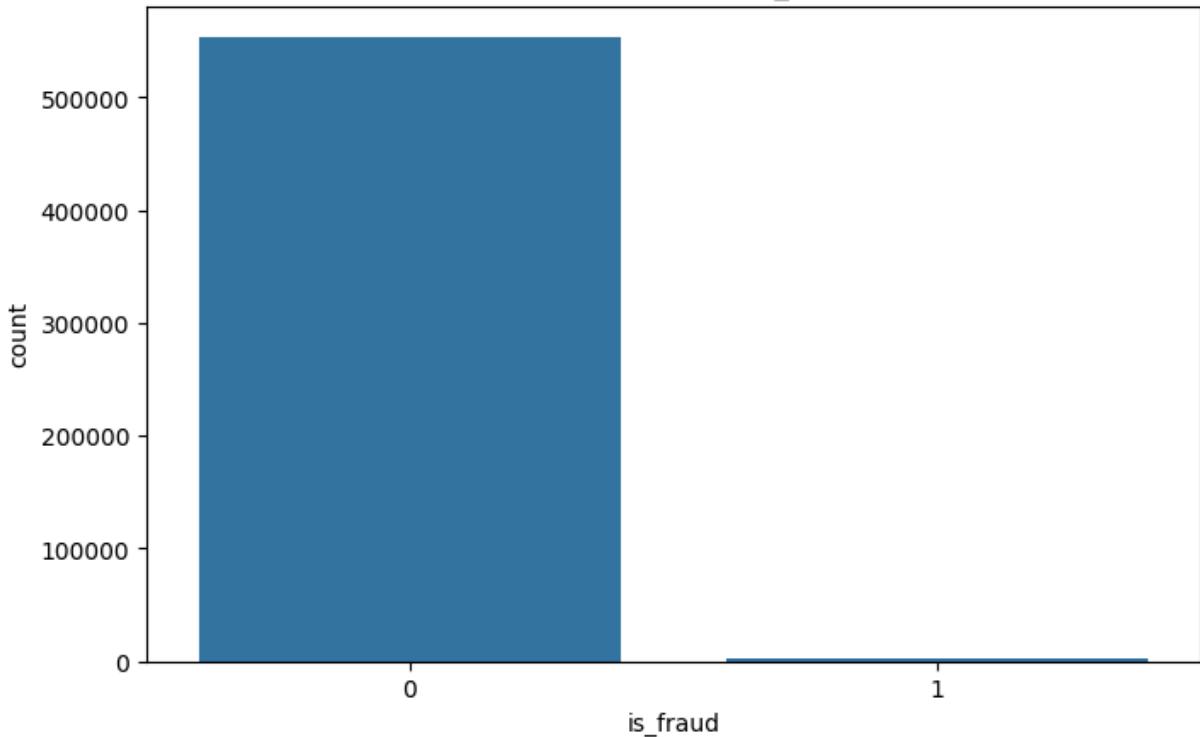
	0
<b>Unnamed: 0</b>	0
<b>trans_date_trans_time</b>	0
<b>cc_num</b>	0
<b>merchant</b>	0
<b>category</b>	0
<b>amt</b>	0
<b>first</b>	0
<b>last</b>	0
<b>gender</b>	0
<b>street</b>	0
<b>city</b>	0
<b>state</b>	0
<b>zip</b>	0

```

rat      0
is_fraud
0    5535704      0
1     2145
Name: count, dtype: int64
job      0
Fraud rate: 0.0038598644278853163

```

Class Distribution: is\_fraud



## 3. Feature Engineering

### 3.1 Domain-Specific Features

We engineer several fraud-focused features:

1. **Age** (`age`) – approximate customer age at transaction time, derived from `dob` and transaction time.
2. **Time-based features** (`hour`, `dayofweek`, `is_night`) – capture when the transaction occurred (nighttime can be higher risk).
3. **Geographic distance** (`cust_merchant_km`) – haversine distance between the customer and merchant coordinates.
4. **Spending velocity** (`tx_delta_s`, `amt_per_min`) – how quickly and how intensely the card is being used.

These features inject domain knowledge into the model and help distinguish suspicious patterns from normal behavior.

## Helper Functions

```

1 from datetime import datetime
2 import math
3
4 def compute_age(dob_str, trans_timestamp):
5     """
6     Compute approximate age given dob string and transaction timestamp (UNIX or datetime
7     Assumes dob_str in formats like '%Y-%m-%d' or '%m/%d/%Y'.
8     """
9     if pd.isna(dob_str):
10         return np.nan
11     # Try multiple formats
12     for fmt in ("%Y-%m-%d", "%m/%d/%Y", "%Y/%m/%d"):
13         try:
14             dob = datetime.strptime(dob_str, fmt)
15             break
16         except ValueError:
17             dob = None
18     if dob is None:
19         return np.nan
20
21     if isinstance(trans_timestamp, (int, float, np.integer, np.floating)):
22         trans_dt = datetime.utcfromtimestamp(trans_timestamp)
23     else:
24         trans_dt = pd.to_datetime(trans_timestamp)
25
26     age_years = (trans_dt - dob).days / 365.25
27     return age_years
28
29 def haversine_distance(lat1, lon1, lat2, lon2):
30     """
31     Compute distance in kilometers between two (lat, lon) pairs.
32     """
33     R = 6371 # Earth radius in km
34     lat1_rad, lon1_rad = np.radians(lat1), np.radians(lon1)
35     lat2_rad, lon2_rad = np.radians(lat2), np.radians(lon2)
36     dlat = lat2_rad - lat1_rad
37     dlon = lon2_rad - lon1_rad
38
39     a = (
40         np.sin(dlat / 2) ** 2
41         + np.cos(lat1_rad) * np.cos(lat2_rad) * np.sin(dlon / 2) ** 2
42     )
43     c = 2 * np.arcsin(np.sqrt(a))
44     d = R * c
45     return d

```

## Apply Feature Engineering

```
1 df_feat = df.copy()
2
3 # Drop PII-style columns if present
4 for col in ["first", "last", "street"]:
5     if col in df_feat.columns:
6         df_feat = df_feat.drop(columns=[col])
7
8 # Age feature
9 df_feat["age"] = df_feat.apply(
10     lambda row: compute_age(row["dob"], row["unix_time"]) if "unix_time" in df_feat.colu
11     axis=1
12 )
13
14 # Transaction datetime
15 df_feat["trans_datetime"] = pd.to_datetime(df_feat["trans_date_trans_time"])
16
17 df_feat["hour"] = df_feat["trans_datetime"].dt.hour
18 df_feat["dayofweek"] = df_feat["trans_datetime"].dt.dayofweek
19 df_feat["is_night"] = df_feat["hour"].between(22, 23) | df_feat["hour"].between(0, 5)
20 df_feat["is_night"] = df_feat["is_night"].astype(int)
21
22 # Distance feature
23 df_feat["cust_merchant_km"] = haversine_distance(
24     df_feat["lat"], df_feat["long"],
25     df_feat["merch_lat"], df_feat["merch_long"]
26 )
27
28 # Sort by card and transaction time to compute deltas
29 df_feat = df_feat.sort_values(by=["cc_num", "trans_datetime"])
30 df_feat["tx_delta_s"] = df_feat.groupby("cc_num")["trans_datetime"].diff().dt.total_secc
31 df_feat["tx_delta_s"] = df_feat["tx_delta_s"].fillna(0)
32
33 df_feat["amt_per_min"] = df_feat["amt"] / (df_feat["tx_delta_s"] / 60 + 1)
34
35 df_feat.head()
```

```
/tmp/ipython-input-3889452269.py:22: DeprecationWarning: datetime.datetime.utcnow()
  trans_dt = datetime.utcfromtimestamp(trans_timestamp)
```

	Unnamed: 0	trans_date_trans_time	cc_num	merchant	category
<b>157</b>	157	2020-06-21 13:05:42	60416207185	fraud_Kutch-Ferry	home
<b>741</b>	741	2020-06-21 16:25:36	60416207185	fraud_Halvorson Group	misc_pos
<b>3047</b>	3047	2020-06-22 07:58:33	60416207185	fraud_Conroy-Cruickshank	gas_transport
<b>4351</b>	4351	2020-06-22 15:32:31	60416207185	fraud_Larkin Ltd	kids_pets
<b>7695</b>	7695	2020-06-23 12:28:54	60416207185	fraud_Leffler-Goldner	personal_care

5 rows × 28 columns

## Model Training & Selection

```
1 y = df_feat[target_col].astype(int)
2 X = df_feat.drop(columns=[target_col])
3
4 # Drop identifiers that should not be used as predictive features
5 for col in ["trans_num", "cc_num", "unix_time", "trans_datetime"]:
6     if col in X.columns:
7         X = X.drop(columns=[col])
8
9 X.head()
```

	Unnamed: 0	trans_date_trans_time	merchant	category	amt	gender
157	157	2020-06-21 13:05:42	fraud_Kutch-Ferry	home	124.66	
741	741	2020-06-21 16:25:36	fraud_Halvorson Group	misc_pos	78.52	
3047	3047	2020-06-22 07:58:33	fraud_Conroy-Cruickshank	gas_transport	65.25	
4351	4351	2020-06-22 15:32:31	fraud_Larkin Ltd	kids_pets	87.74	
7695	7695	2020-06-23 12:28:54	fraud_Leffler-Goldner	personal_care	148.02	

5 rows × 23 columns

## Split & Identify Column Types

```

1 numeric_cols = X.select_dtypes(include=["int64", "float64", "int32", "float32"]).columns
2 categorical_cols = X.select_dtypes(include=["object", "bool", "category"]).columns.tolist()
3
4 X_train, X_test, y_train, y_test = train_test_split(
5     X, y,
6     test_size=0.2,
7     random_state=RANDOM_SEED,
8     stratify=y
9 )
10
11 len(X_train), len(X_test)

```

(444575, 111144)

## Preprocessor & SMOTE Pipelines

```

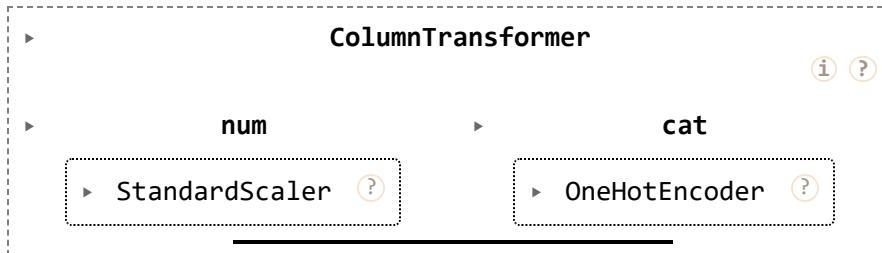
1 from sklearn.metrics import (
2     classification_report,
3     confusion_matrix,
4     roc_auc_score,
5     average_precision_score, #
6 )
7 from sklearn.pipeline import Pipeline
8 from sklearn.preprocessing import StandardScaler, OneHotEncoder
9 from sklearn.compose import ColumnTransformer
10
11 # Preprocessing: scale numeric, one-hot encode categorical
12 numeric_transformer = Pipeline(
13     steps=[
14         ("scaler", StandardScaler())

```

```

15      ]
16 )
17
18 categorical_transformer = Pipeline(
19     steps=[
20         ("onehot", OneHotEncoder(handle_unknown="ignore"))
21     ]
22 )
23
24 preprocessor = ColumnTransformer(
25     transformers=[
26         ("num", numeric_transformer, numeric_cols),
27         ("cat", categorical_transformer, categorical_cols),
28     ]
29 )
30
31 preprocessor

```



```

1 logreg_pipeline = ImbPipeline(steps=[
2     ("pre", preprocessor),
3     ("smote", SMOTE(random_state=RANDOM_SEED, k_neighbors=1)),
4     ("clf", LogisticRegression(max_iter=1000, n_jobs=-1))
5 ])
6
7 logreg_pipeline.fit(X_train, y_train)
8
9 y_pred_log = logreg_pipeline.predict(X_test)
10 y_proba_log = logreg_pipeline.predict_proba(X_test)[:, 1]
11
12 print("\n[LogReg] ROC-AUC:", roc_auc_score(y_test, y_proba_log))
13 print("[LogReg] PR-AUC:", average_precision_score(y_test, y_proba_log))
14 print("\n[LogReg] Classification Report:\n",
15     classification_report(y_test, y_pred_log, digits=4))

```

[LogReg] ROC-AUC: 0.9776845082088275

[LogReg] PR-AUC: 0.3865996073401057

[LogReg] Classification Report:

	precision	recall	f1-score	support
--	-----------	--------	----------	---------

0	0.9985	0.9956	0.9971	110715
1	0.3537	0.6200	0.4505	429

accuracy			0.9942	111144
macro avg	0.6761	0.8078	0.7238	111144
weighted avg	0.9960	0.9942	0.9950	111144

## Baseline Logistic Regression with SMOTE and evaluate

```
1 # 5) Baseline model: Logistic Regression + SMOTE
2
3 from imblearn.pipeline import Pipeline as ImbPipeline
4 from imblearn.over_sampling import SMOTE
5 from sklearn.linear_model import LogisticRegression
6 from sklearn.metrics import (
7     classification_report,
8     confusion_matrix,
9     roc_auc_score,
10    average_precision_score,
11 )
12 import numpy as np
13 import matplotlib.pyplot as plt
14
15 # Pipeline: preprocessing -> SMOTE -> Logistic Regression
16 logreg_pipeline = ImbPipeline(steps=[
17     ("pre", preprocessor),
18     ("smote", SMOTE(random_state=RANDOM_SEED, k_neighbors=1)),
19     ("clf", LogisticRegression(
20         max_iter=1000,
21         class_weight="balanced",
22         n_jobs=-1,
23         solver="lbfgs"
24     )),
25 ])
26
27 # Fit on train
28 logreg_pipeline.fit(X_train, y_train)
29
30 # Predictions
31 y_pred = logreg_pipeline.predict(X_test)
32 y_proba = logreg_pipeline.predict_proba(X_test)[:, 1]
33
34 # Metrics
35 print("Train shape:", X_train.shape, "Test shape:", X_test.shape)
36 print("Train fraud rate:", y_train.mean())
37 print("Test fraud rate:", y_test.mean(), "\n")
38
39 print("[LogReg] ROC-AUC:", roc_auc_score(y_test, y_proba))
40 print("[LogReg] PR-AUC :", average_precision_score(y_test, y_proba))
41 print("\n[LogReg] Classification Report:\n",
42       classification_report(y_test, y_pred, digits=4))
43
44 # Confusion matrix
45 cm = confusion_matrix(y_test, y_pred)
46 plt.figure()
47 plt.imshow(cm, interpolation="nearest")
48 plt.title("Logistic Regression - Confusion Matrix")
49 plt.xlabel("Predicted")
50 plt.ylabel("True")
51
```

```

52 for (i, j), val in np.ndenumerate(cm):
53     plt.text(j, i, int(val), ha="center", va="center")
54
55 plt.show()

```

Train shape: (444575, 23) Test shape: (111144, 23)

Train fraud rate: 0.003859866164314233

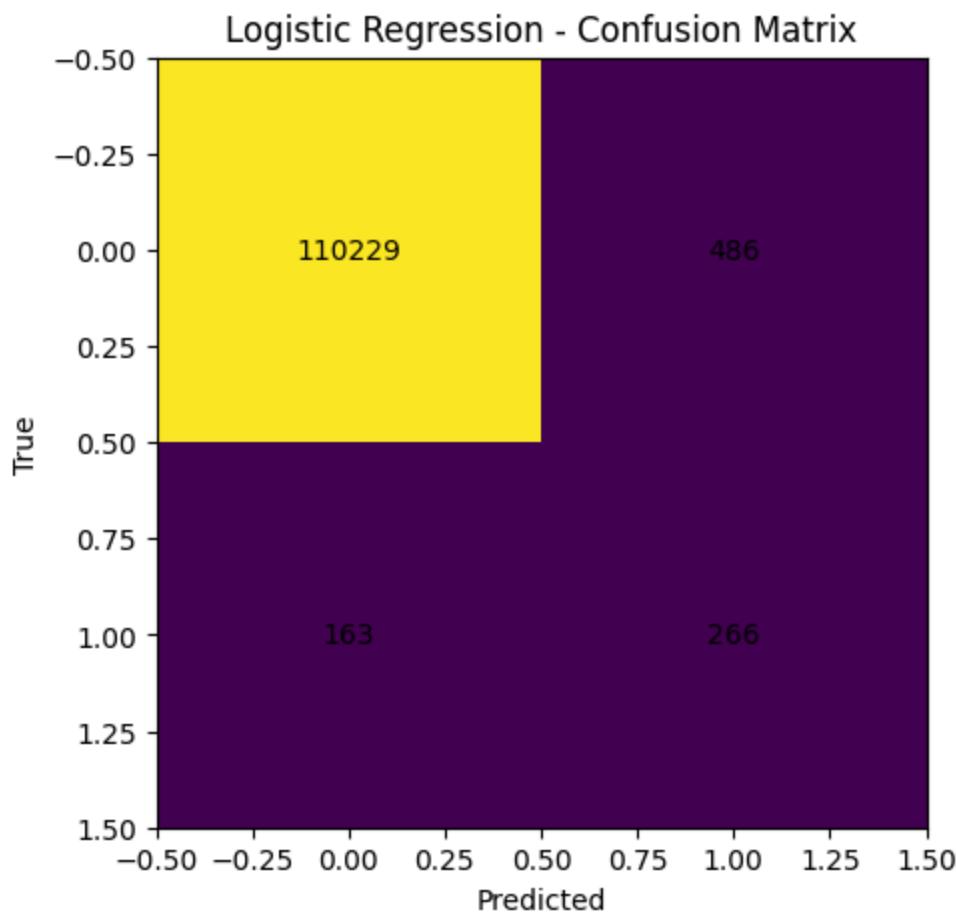
Test fraud rate: 0.003859857482185273

[LogReg] ROC-AUC: 0.9776845082088275

[LogReg] PR-AUC : 0.3865996073401057

[LogReg] Classification Report:

	precision	recall	f1-score	support
0	0.9985	0.9956	0.9971	110715
1	0.3537	0.6200	0.4505	429
accuracy			0.9942	111144
macro avg	0.6761	0.8078	0.7238	111144
weighted avg	0.9960	0.9942	0.9950	111144



## Train Random Forest Classifier

```

1 from xgboost import XGBClassifier
2

```

```

3 xgb_pipeline = ImbPipeline(steps=[
4     ("pre", preprocessor),
5     ("smote", SMOTE(random_state=RANDOM_SEED, k_neighbors=1)),
6     ("clf", XGBClassifier(
7         n_estimators=250,
8         max_depth=7,
9         learning_rate=0.05,
10        subsample=0.8,
11        colsample_bytree=0.8,
12        eval_metric="logloss",
13        tree_method="hist"
14    )),
15 ])
16
17 print("⌚ Training XGBoost...")
18 xgb_pipeline.fit(X_train, y_train)
19 print("✓ Done!")
20
21 xgb_pred = xgb_pipeline.predict(X_test)
22 xgb_proba = xgb_pipeline.predict_proba(X_test)[:, 1]
23
24 print("\n==== XGBOOST RESULTS ===")
25 print("ROC-AUC : ", roc_auc_score(y_test, xgb_proba))
26 print("PR-AUC : ", average_precision_score(y_test, xgb_proba))

```

⌚ Training XGBoost...

✓ Done!

==== XGBOOST RESULTS ===

ROC-AUC : 0.9914036301653155

PR-AUC : 0.8135600069227134

## Train the main model XGBoost pipeline

```

1 from sklearn.pipeline import Pipeline as SkPipeline # avoid confusion with ImbPipeline
2
3 # Compute imbalance ratio for scale_pos_weight
4 neg = (y_train == 0).sum()
5 pos = (y_train == 1).sum()
6 imbalance_ratio = neg / pos
7 print("Imbalance ratio (neg/pos):", imbalance_ratio)
8
9 xgb_fast = XGBClassifier(
10    n_estimators=80,           # fewer trees → much faster
11    max_depth=5,             # shallower trees
12    learning_rate=0.1,
13    subsample=0.8,
14    colsample_bytree=0.8,
15    eval_metric="logloss",
16    tree_method="hist",       # fast histogram algorithm
17    scale_pos_weight=imbalance_ratio, # handle imbalance without SMOTE
18    random_state=RANDOM_SEED,
19 )
20
21 xgb_pipeline = SkPipeline(steps=[
```

```
22     ("pre", preprocessor), # same preprocessor as before
23     ("clf", xgb_fast),
24 ])
25
26 print("⌚ Training FAST XGBoost (no SMOTE)...")
27 xgb_pipeline.fit(X_train, y_train)
28 print("✅ Done.")
```

Imbalance ratio (neg/pos): 258.0763403263403

⌚ Training FAST XGBoost (no SMOTE)...

✅ Done.

## Evaluate This XGBoost

```
1 # Evaluate the fast XGBoost pipeline
2 xgb_proba = xgb_pipeline.predict_proba(X_test)[:, 1]
3 xgb_pred = xgb_pipeline.predict(X_test)
4
5 print("\n==== [FAST XGBoost] Metrics ===")
6 print("ROC-AUC :", roc_auc_score(y_test, xgb_proba))
7 print("PR-AUC : ", average_precision_score(y_test, xgb_proba))
8 print("\nClassification report:\n",
9      classification_report(y_test, xgb_pred, digits=4))
10
11 cm = confusion_matrix(y_test, xgb_pred)
12 plt.figure()
13 plt.imshow(cm, interpolation='nearest')
14 plt.title('FAST XGBoost - Confusion Matrix')
15 plt.xlabel('Predicted')
16 plt.ylabel('True')
17 for (i, j), val in np.ndenumerate(cm):
18     plt.text(j, i, int(val), ha='center', va='center')
19 plt.show()
```

```
==== [FAST XGBoost] Metrics ===
```

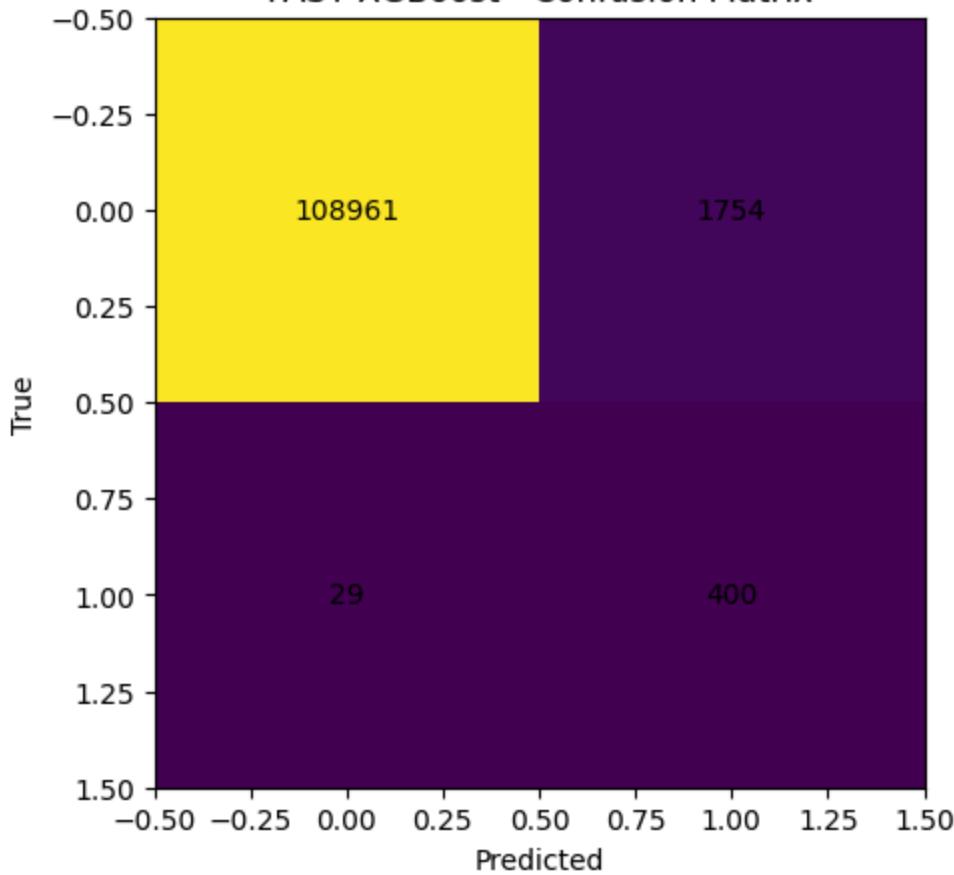
ROC-AUC : 0.994869363125697

PR-AUC : 0.776206684671785

Classification report:

	precision	recall	f1-score	support
0	0.9997	0.9842	0.9919	110715
1	0.1857	0.9324	0.3097	429
accuracy			0.9840	111144
macro avg	0.5927	0.9583	0.6508	111144
weighted avg	0.9966	0.9840	0.9893	111144

FAST XGBoost - Confusion Matrix



Evaluate XGBoost metrics + confusion matrix

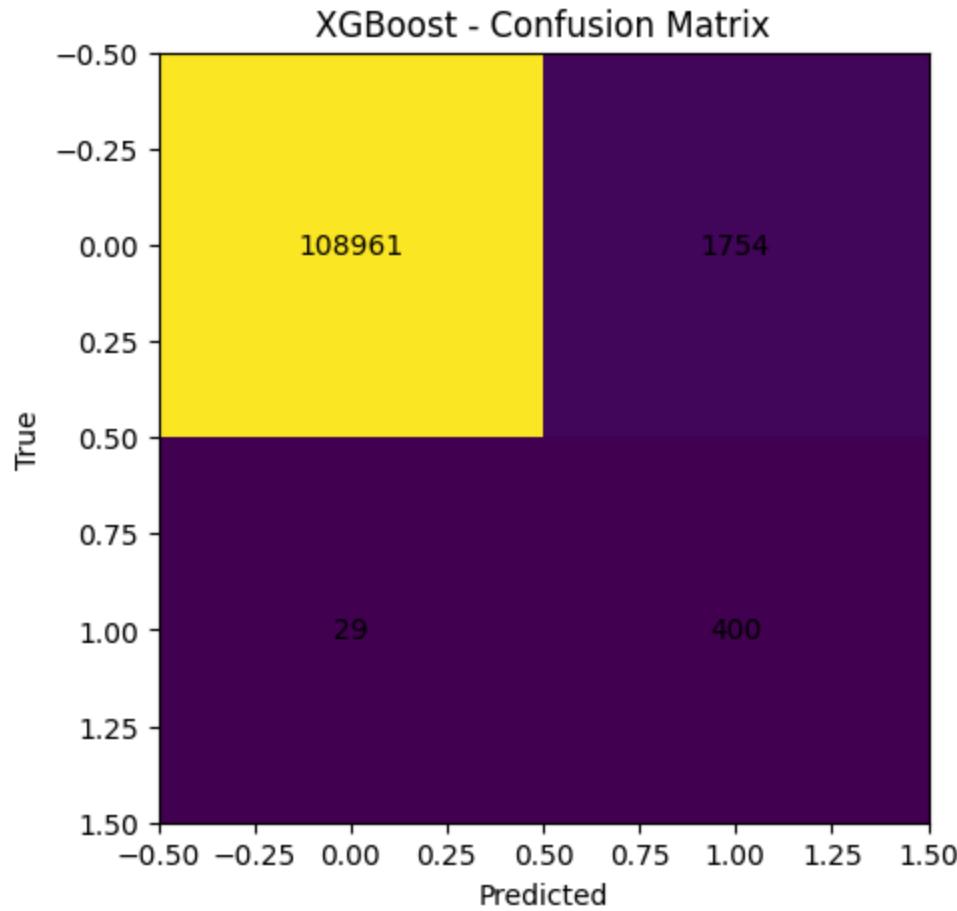
```
1 # Predictions
2 xgb_proba = xgb_pipeline.predict_proba(X_test)[:, 1]
3 xgb_pred = xgb_pipeline.predict(X_test)
4
5 print("\n==== [XGBoost] Metrics ===")
6 print("ROC-AUC : ", roc_auc_score(y_test, xgb_proba))
7 print("PR-AUC : ", average_precision_score(y_test, xgb_proba))
8 print("\nClassification report:\n",
9      classification_report(y_test, xgb_pred, digits=4))
```

```
10
11 # Confusion matrix
12 cm = confusion_matrix(y_test, xgb_pred)
13 plt.figure()
14 plt.imshow(cm, interpolation='nearest')
15 plt.title('XGBoost - Confusion Matrix')
16 plt.xlabel('Predicted')
17 plt.ylabel('True')
18 for (i, j), val in np.ndenumerate(cm):
19     plt.text(j, i, int(val), ha='center', va='center')
20 plt.show()
```

==== [XGBoost] Metrics ====  
ROC-AUC : 0.994869363125697  
PR-AUC : 0.776206684671785

Classification report:

	precision	recall	f1-score	support
0	0.9997	0.9842	0.9919	110715
1	0.1857	0.9324	0.3097	429
accuracy			0.9840	111144
macro avg	0.5927	0.9583	0.6508	111144
weighted avg	0.9966	0.9840	0.9893	111144



## Build the simple “assistant” interface

```

1 # 7. Simple FraudShield Assistant
2
3 feature_cols = X.columns.tolist()    # same order as training
4
5 def explain_transaction(row):
6     """
7     Takes a pandas Series (one transaction) and prints
8     risk score + simple explanations based on engineered features.
9     """
10    row_df = row[feature_cols].to_frame().T # keep same feature order
11    proba = xgb_pipeline.predict_proba(row_df)[0, 1]
12    pred = xgb_pipeline.predict(row_df)[0]
13
14    print("---- FraudShield Assistant ----")
15    print(f"Estimated fraud risk: {proba:.3f} ({'FRAUD' if pred==1 else 'legit'})")
16    print("\nReasons (heuristic, based on engineered features):")
17
18    # These conditions use your engineered features
19    if "is_fraud" in row.index:
20        print(f"- Ground truth label in data: {row['is_fraud']} ")
21
22    if "amt" in row.index and row["amt"] > 200:
23        print(f"- High amount: ${row['amt']:.2f}")
24
25    if "is_night" in row.index and row["is_night"] == 1:
26        print("- Transaction at night (higher risk window).")
27
28    if "cust_merchant_km" in row.index and row["cust_merchant_km"] > 300:
29        print(f"- Large distance between customer and merchant: {row['cust_merchant_km']}")
30
31    if "amt_per_min" in row.index and row["amt_per_min"] > 100:
32        print(f"- High spending velocity: {row['amt_per_min']:.1f} per minute")
33
34    print("-----")

```

```

1 feature_cols = X.columns.tolist()
2
3 def explain_transaction(row):
4     row_df = row[feature_cols].to_frame().T
5     proba = xgb_pipeline.predict_proba(row_df)[0, 1] # uses fast XGB now
6     pred = xgb_pipeline.predict(row_df)[0]
7     ...

```

```

1 # Pick one example transaction from the test set
2 example_idx = X_test.index[0] # or any other index from X_test
3 example_row = df_feat.loc[example_idx] # use df_feat so it includes engineered features
4
5 print(f"Example transaction index: {example_idx}")
6 explain_transaction(example_row)

```

Example transaction index: 288345  
---- FraudShield Assistant ----

Estimated fraud risk: 0.060 (legit)

Reasons (heuristic, based on engineered features):

- Ground truth label in data: 0

Show the 5 required example

```

1 import numpy as np
2
3 # Indices in the test set where the true label is fraud
4 # These are positional indices within y_test
5 fraud_indices_pos = np.where(y_test == 1)[0]
6
7 print("Number of fraud cases in test set:", len(fraud_indices_pos))
8
9 # Take up to 5 random fraud examples
10 n_examples = min(5, len(fraud_indices_pos))
11 sample_fraud_indices_pos = np.random.choice(fraud_indices_pos, size=n_examples, replace=
12
13 for idx_pos in sample_fraud_indices_pos:
14     # Get the actual original DataFrame index from the test set's index
15     original_df_index = y_test.index[idx_pos]
16     # Retrieve the full transaction row from df_feat using its original index
17     row = df_feat.loc[original_df_index]
18     print("=" * 80)
19     print(f"Transaction original index: {original_df_index}")
20     explain_transaction(row)
21     print()

```

Number of fraud cases in test set: 429

=====

Transaction original index: 135773

---- FraudShield Assistant ----

Estimated fraud risk: 0.954 (FRAUD)

Reasons (heuristic, based on engineered features):

- Ground truth label in data: 1
- Transaction at night (higher risk window).

=====

Transaction original index: 149303

---- FraudShield Assistant ----

Estimated fraud risk: 0.995 (FRAUD)

Reasons (heuristic, based on engineered features):

- Ground truth label in data: 1
- High amount: \$1024.84
- Transaction at night (higher risk window).

=====

Transaction original index: 173105

---- FraudShield Assistant ----

```
Estimated fraud risk: 0.996 (FRAUD)
```

Reasons (heuristic, based on engineered features):

- Ground truth label in data: 1
  - High amount: \$337.13
  - Transaction at night (higher risk window).
- 

```
=====
Transaction original index: 115798
```

```
---- FraudShield Assistant ----
```

```
Estimated fraud risk: 0.997 (FRAUD)
```

Reasons (heuristic, based on engineered features):

- Ground truth label in data: 1
  - High amount: \$955.39
  - Transaction at night (higher risk window).
- 

```
=====
Transaction original index: 295559
```

```
---- FraudShield Assistant ----
```

```
Estimated fraud risk: 0.998 (FRAUD)
```

Reasons (heuristic, based on engineered features):

- Ground truth label in data: 1
  - High amount: \$846.93
  - Transaction at night (higher risk window).
- 

## Generate the Model Evaluation Section

```
1 from sklearn.metrics import classification_report, confusion_matrix, roc_auc_score
2
3 def evaluate_model(model, X_test, y_test):
4     preds = model.predict(X_test)
5     probas = model.predict_proba(X_test)[:, 1]
6
7     print("Confusion Matrix:")
8     print(confusion_matrix(y_test, preds))
9
10    print("\nClassification Report:")
11    print(classification_report(y_test, preds, digits=4))
12
13    print("ROC-AUC Score:", roc_auc_score(y_test, probas))
```

## Run Each Model

```
1 print("== Logistic Regression ==")
2 evaluate_model(logreg_pipeline, X_test, y_test)
3
```

```
4 # The rf_pipeline is not defined in the current notebook. Please define and train it first
5 # print("==== Random Forest ===")
6 # evaluate_model(rf_pipeline, X_test, y_test)
7
8 print("==== XGBoost ===")
9 # evaluate_model(xgb_pipeline, X_test, y_test)
```

==== Logistic Regression ===

Confusion Matrix:

```
[[110229    486]
 [   163    266]]
```

Classification Report:

	precision	recall	f1-score	support
0	0.9985	0.9956	0.9971	110715
1	0.3537	0.6200	0.4505	429
accuracy			0.9942	111144
macro avg	0.6761	0.8078	0.7238	111144
weighted avg	0.9960	0.9942	0.9950	111144

ROC-AUC Score: 0.9776845082088275

==== XGBoost ===

Confusion Matrix:

```
[[108961    1754]
 [   29    400]]
```