

✓ Title and Introduction

Data Preprocessing Lab: Generative AI

Welcome to the Data Preprocessing Lab for Generative AI!

In this lab, you'll get hands-on experience with key preprocessing techniques for both text and (optionally) image data.

Learning Objectives:

- Understand and apply core data preprocessing techniques.
- Explore word embedding techniques (Word2Vec/GloVe, BERT).
- Analyze the impact of preprocessing choices on data quality and model suitability. List item
- Practice using cosine similarity for comparing embeddings.

✓ Part 1: Environment Setup

First, we'll install and import all necessary libraries. Run the following cell to set up your environment.

```
# SECTION 1: Environment Setup
#####
# This cell installs and imports all necessary libraries for our text preprocessing pipeline.
# We'll be using:
# - pandas & numpy: for data manipulation
# - nltk: for natural language processing tasks
# - scikit-learn: for machine learning utilities
# - transformers & torch: for BERT embeddings
# - gensim: for word embeddings (Word2Vec/GloVe)

# Install required packages
%pip uninstall -y numpy pandas
%pip install pandas numpy==1.26.4 nltk scikit-learn transformers torch datasets gensim

# TODO: Import the required libraries
# Hint: You need pandas, numpy, nltk, and sklearn components
# YOUR CODE HERE - import the basic libraries
import pandas as pd
import numpy as np
import nltk
# Add more imports as needed...

# These are more advanced imports you'll need later
from transformers import BertTokenizer, BertModel
import torch
import gensim.downloader as api
from gensim.models import KeyedVectors
from sklearn.model_selection import train_test_split
from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.metrics.pairwise import cosine_similarity

# Download required NLTK data
# These are necessary for tokenization, stop words, and lemmatization
nltk.download('punkt') # Added punkt_tab
nltk.download('stopwords')
nltk.download('wordnet')
nltk.download('omw-1.4')

print("Setup complete! All required libraries have been imported.")
```



```

Found existing installation: numpy 1.26.4
Uninstalling numpy-1.26.4:
  Successfully uninstalled numpy-1.26.4
Found existing installation: pandas 2.3.3
Uninstalling pandas-2.3.3:
  Successfully uninstalled pandas-2.3.3
Collecting pandas
  Using cached pandas-2.3.3-cp312-cp312-manylinux_2_24_x86_64.manylinux_2_28_x86_64.whl.metadata (91 kB)
Collecting numpy==1.26.4
  Using cached numpy-1.26.4-cp312-cp312-manylinux_2_17_x86_64.manylinux2014_x86_64.whl.metadata (61 kB)
Requirement already satisfied: nltk in /usr/local/lib/python3.12/dist-packages (3.9.1)
Requirement already satisfied: scikit-learn in /usr/local/lib/python3.12/dist-packages (1.6.1)
Requirement already satisfied: transformers in /usr/local/lib/python3.12/dist-packages (4.57.0)
Requirement already satisfied: torch in /usr/local/lib/python3.12/dist-packages (2.8.0+cu126)
Requirement already satisfied: datasets in /usr/local/lib/python3.12/dist-packages (4.0.0)
Requirement already satisfied: gensim in /usr/local/lib/python3.12/dist-packages (4.3.3)
Requirement already satisfied: python-dateutil>=2.8.2 in /usr/local/lib/python3.12/dist-packages (from pandas) (2.9.0.post0)
Requirement already satisfied: pytz>=2020.1 in /usr/local/lib/python3.12/dist-packages (from pandas) (2025.2)
Requirement already satisfied: tzdata>=2022.7 in /usr/local/lib/python3.12/dist-packages (from pandas) (2025.2)
Requirement already satisfied: click in /usr/local/lib/python3.12/dist-packages (from nltk) (8.3.0)
Requirement already satisfied: joblib in /usr/local/lib/python3.12/dist-packages (from nltk) (1.5.2)
Requirement already satisfied: regex>=2021.8.3 in /usr/local/lib/python3.12/dist-packages (from nltk) (2024.11.6)
Requirement already satisfied: tqdm in /usr/local/lib/python3.12/dist-packages (from nltk) (4.67.1)
Requirement already satisfied: scipy>=1.6.0 in /usr/local/lib/python3.12/dist-packages (from scikit-learn) (1.13.1)
Requirement already satisfied: threadpoolctl>=3.1.0 in /usr/local/lib/python3.12/dist-packages (from scikit-learn) (3.6.0)
Requirement already satisfied: filelock in /usr/local/lib/python3.12/dist-packages (from transformers) (3.20.0)
Requirement already satisfied: requests in /usr/local/lib/python3.12/dist-packages (from transformers) (0.35.3)
Requirement already satisfied: packaging>=20.0 in /usr/local/lib/python3.12/dist-packages (from transformers) (25.0)
Requirement already satisfied: pyyaml>=5.1 in /usr/local/lib/python3.12/dist-packages (from transformers) (6.0.3)
Requirement already satisfied: tokenizers<=0.23.0,>=0.22.0 in /usr/local/lib/python3.12/dist-packages (from transformers) (0.22.1)

```

Part 2: Loading and Exploring the BBC News Dataset

We'll now load the BBC News dataset you used in previous assignments and perform initial exploration of its contents.

```
Requirement already satisfied: tokenizers<=0.23.0,>=0.22.0 in /usr/local/lib/python3.12/dist-packages (from transformers) (0.22.1)
```

```

# SECTION 2: Data Loading and Initial Exploration
#####
# Here we load the BBC News dataset and perform initial analysis
# Understanding our data is crucial before applying any preprocessing

# Load the dataset
from google.colab import files
import pandas as pd

# Use the uploaded filename exactly
# Assuming the uploaded file is named 'bbc-news-data.csv'
# You might need to adjust the filename if it's different after uploading
try:
    uploaded = files.upload()
    file_name = list(uploaded.keys())[0]
    df = pd.read_csv(file_name, encoding='latin1')

    # Verify load
    print("Dataset loaded successfully with shape:", df.shape)
    display(df.head())

    # TODO: Rename the columns to match our processing pipeline
    # Hint: The original columns are 'Text' and 'Category'
    # YOUR CODE HERE
    # Assuming the columns in the CSV are 'labels' and 'text' based on the kernel state
    df.rename(columns={'labels': 'category', 'text': 'text'}, inplace=True)
    print("Renamed columns to 'category' and 'text'")

    # Verify that the 'category' column exists after renaming
    if 'category' not in df.columns:
        raise KeyError("Column 'category' not found after renaming. Please check the original column names in your CSV file.")

    # TODO: Perform basic data exploration
    # TASK 1: Display the first few rows and basic information about the dataset
    # Hint: Use pandas' head(), info(), and describe() methods
    # YOUR CODE HERE
    print("\nDataset Info:")
    df.info()
    print("\nDataset Description:")
    display(df.describe(include='all'))

    # TASK 2: Analyze the distribution of categories
    # Hint: Use value_counts() on the category column
    # YOUR CODE HERE
    print("\nCategory Distribution:")
    display(df['category'].value_counts())

    # TASK 3: Calculate and display basic text statistics
    # Calculate average text length per category
    # TODO: Create a visualization of text lengths by category
    # Hint: Use seaborn's boxplot

```

```
# YOUR CODE HERE
df['text_length'] = df['text'].str.len()
print("\nText Length Statistics per Category:")
display(df.groupby('category')['text_length'].describe())

# TODO: Create a visualization of text lengths by category
# Hint: Use seaborn's boxplot
# YOUR CODE HERE
import seaborn as sns
import matplotlib.pyplot as plt

plt.figure(figsize=(10, 6))
sns.boxplot(x='category', y='text_length', data=df)
plt.title('Distribution of Text Lengths by Category')
plt.xlabel('Category')
plt.ylabel('Text Length')
plt.xticks(rotation=45)
plt.show()

# Display your findings
print("\nDataset Statistics:")
# TODO: Add code to display your findings
# YOUR CODE HERE
# Already displayed above with describe(), value_counts(), and groupby().describe()

except FileNotFoundError:
    print("Please upload the 'bbc-news-data.csv' file using the files.upload() function.")
except KeyError as e:
    print(f"KeyError: {e}. Please check the column names in your CSV file.")
except Exception as e:
    print(f"An unexpected error occurred: {e}")
```

Part 3: Text Preprocessing

We'll now implement basic text preprocessing steps to clean our data.

▼ Comprehension Questions - Data Exploration

Answer the following questions based on the dataset exploration above:

1. What are the dimensions of our dataset? The datasets has 7 columns and 100 rows.
2. How many different categories are there in the news articles? We have 7 categories which are: text, labels, no_sentences, Flesch Reading Ease Score, Dale-Chall Readability Score, text_rank_summary, lsa_summary

3. Is the dataset balanced across categories? Why might this matter? The Data is balanced because it has the same numbers of samples. 100 rows dataset balanced 27% left by the imbalance to 100% done
Saving bbc news.100rows.(1).csv to bbc news.100rows.(1).(9).csv
dataset loaded successfully with shape: (100, 7)
4. Are there any missing values that need to be addressed? This csv file has been modified since it has 1000 rows and I had to cut it down to 100 rows and there is no missing values in the new csv file.

Flesch **Dale-Chall** **text_rank_summary** **lca_summary**

```
# SECTION 3: Text Cleaning and Preprocessing
#####
# This section implements fundamental text preprocessing steps:
# 1. Converting to lowercase (why? -> maintains consistency)
# 2. Removing special characters (why? -> reduces noise)
# 3. Handling whitespace (why? -> standardizes format)

#####
import re # Moved import to the top

def clean_text(text):
    """
    Performs basic text cleaning operations.

    Parameters:
    text (str): Input text to be cleaned

    Returns:
    str: Cleaned text
    """

    # TODO: Implement the following steps:
    # 1. Convert to lowercase
    # 2. Remove URLs and emails
    # 3. Remove special characters but keep sentence structure
    # 4. Remove extra whitespace
    # Hint: Use string methods and regular expressions

    # YOUR CODE HERE
    text = str(text).lower() # 1) lowercase
    text = re.sub(r"(https?:\/\/S+|www.\S+)", ' ', text) # 2) remove URLs
    text = re.sub(r'\b[\w\.-]+@[^\w\.-]+\.\w+\b', ' ', text) # 2) remove emails
    text = re.sub(r"[^a-z0-9\.,\!?\.:;\\""]", " ", text) # 3) remove special chars but keep basic sentence punctuation
    # normalize spacing before punctuation (moved inside the function)
    text = re.sub(r"\s+(.,!?:;)", r"\1", text)
    # 4) collapse extra whitespace
    text = re.sub(r"\s+", " ", text).strip()

    return text

# Test the function with a sample (Corrected indentation)
sample_text = "Hello, World! This is a TEST... 123. Check out this link: https://example.com and email: test@example.com"
print("Original:", sample_text)
print("Cleaned:", clean_text(sample_text))

# Apply to the entire dataset (Corrected indentation)
df['cleaned_text'] = df['text'].apply(clean_text)
```

unique	100	1	NaN	NaN	NaN	NaN	100	100
Original:	Hello, World! This is a TEST... 123. Check out this link: https://example.com and email: test@example.com							
Cleaned:	hello world this is a test... 123. check out this link: and email:							
top	Warner	business	NaN	NaN	NaN	It hopes to increase	subscribers by offering t...	Its profits were
	profit\n\nQuarterly...							buoyed by one-off
								gains which...

Part 4: Tokenization and Advanced Processing

mean	100	1	NaN	NaN	NaN	NaN	100	100
Now we'll tokenize our text and apply more advanced preprocessing techniques including:								
std	NaN	NaN	16.730000	61.236700	9.625900	NaN	NaN	NaN
• Tokenization	NaN	NaN	5.136235	7.903453	0.606077	NaN	NaN	NaN
• Stop word removal	NaN	NaN	7.000000	34.360000	8.340000	NaN	NaN	NaN
• Lemmatization	NaN	NaN	13.000000	55.662500	9.167500	NaN	NaN	NaN
50%	NaN	NaN	16.000000	61.415000	9.665000	NaN	NaN	NaN

```
# SECTION 4: Tokenization and Advanced Processing
#####
# This section implements more sophisticated NLP techniques:
# - Tokenization: splitting text into words
# - Stop word removal: removing common words
# - Lemmatization: reducing words to their base form
# Check if you do not need to install any additional libraries
from nltk.corpus import stopwords
from nltk.tokenize import word_tokenize
from nltk.stem import WordNetLemmatizer

# Initialize our tools
stop_words = set(stopwords.words('english'))
lemmatizer = WordNetLemmatizer()

def tokenize_and_process(text):
```

```
"""
Performs advanced text processing including tokenization,
stop word removal, and lemmatization.

Parameters:
text (str): Cleaned text to process

Returns:
list: List of processed tokens
"""

# TODO: Implement the following steps:
# 1. Tokenize the text
# 2. Remove stop words
# 3. Apply lemmatization
# Hint: Use the initialized stop_words and lemmatizer

# YOUR CODE HERE
tokens = word_tokenize(text) # 1. Tokenize the text
processed_tokens = [
    lemmatizer.lemmatize(word) for word in tokens if word not in
stop_words # 2. Remove stop words and 3. Apply lemmatization
]
return processed_tokens

# Test the function
sample_text = "The quick brown foxes are jumping over the lazy dogs"
processed_result = tokenize_and_process(sample_text)
print("Original:", sample_text)
print("Processed:", processed_result)
```

Original: The quick brown foxes are jumping over the lazy dogs
 Processed: ['The', 'quick', 'brown', 'fox', 'jumping', 'lazy', 'dog']

Part 5: Word Embeddings with GloVe

We'll now generate word embeddings using pre-trained GloVe vectors. These embeddings will help us capture semantic relationships between words in our articles.

```
# SECTION 5: Word Embeddings with GloVe
#####
# This section generates word embeddings using pre-trained GloVe vectors
# Word embeddings capture semantic relationships between words
# by representing them as dense vectors in a high-dimensional space

import gensim.downloader as api
import numpy as np # Import numpy
from nltk.tokenize import word_tokenize # Import word_tokenize
from nltk.stem import WordNetLemmatizer # Import WordNetLemmatizer
from nltk.corpus import stopwords # Import stopwords

# Initialize NLTK tools needed for tokenization and processing within the embedding function
stop_words = set(stopwords.words('english'))
lemmatizer = WordNetLemmatizer()

# Load pre-trained GloVe embeddings
glove_model = api.load("glove-wiki-gigaword-100")

def tokenize_and_process_for_glove(text):
    """
    Performs basic text processing including tokenization,
    stop word removal, and lemmatization suitable for GloVe.

    Parameters:
    text (str): Cleaned text to process

    Returns:
    list: List of processed tokens
    """

    # 1) Tokenize
    tokens = word_tokenize(str(text))

    # 2) Remove stop words (and drop non-alphabetic tokens); lowercase for consistency
    tokens = [t.lower() for t in tokens if t.isalpha() and t.lower() not in stop_words]

    # 3) Lemmatize
    processed_tokens = [lemmatizer.lemmatize(t) for t in tokens]

    return processed_tokens

def get_word2vec_embedding(text, model):
```

```
"""
Generates document embeddings by averaging word vectors.

Parameters:
text (str): Input text (cleaned text is expected)
model: Pre-trained word embedding model

Returns:
numpy.array: Document embedding vector
"""

# TODO: Implement the following steps:
# 1. Tokenize the input text
# 2. Get embedding for each token
# 3. Average the embeddings
# Hint: Handle words not in vocabulary

# YOUR CODE HERE
# 1) Tokenize and process the text using the helper function
processed_tokens = tokenize_and_process_for_glove(text)

# 2) collect vectors for in-vocab tokens (case-insensitive fallback)
vecs = []
for w in processed_tokens:
    if w in model.key_to_index:
        vecs.append(model[w])
    elif w.lower() in model.key_to_index:
        vecs.append(model[w.lower()])

# 3) Average the embeddings
return np.mean(vecs, axis=0) if vecs else np.zeros(model.vector_size)

# Apply to a sample of the dataset
sample_size = 100
sample_df = df.head(sample_size).copy()
sample_df['glove_embedding'] = sample_df['cleaned_text'].apply(
    lambda x: get_word2vec_embedding(x, glove_model)
)
print("Sample of embeddings:")
# Apply to a sample of the dataset
sample_size = 100
[=====] 100.0%
128.1/128.1MB downloaded
```

Sample of embeddings:

Part 6: BERT Embeddings

Now we'll use BERT to generate contextual embeddings. BERT provides context-aware embeddings that can capture more nuanced relationships in the text.

```
# SECTION 6: BERT Embeddings
#####
# This section implements BERT (Bidirectional Encoder Representations from Transformers)
# BERT provides context-aware embeddings, meaning the same word can have different
# embeddings based on its context in the sentence.
# Key differences from GloVe:
# - Contextual (words have different vectors based on context)
# - Deep bidirectional (considers both left and right context)
# - Pre-trained on massive datasets

from transformers import BertTokenizer, BertModel # Added import
import torch # Added import
import numpy as np # Added import for numpy usage

# Load BERT model and tokenizer
tokenizer = BertTokenizer.from_pretrained('bert-base-uncased')
model = BertModel.from_pretrained('bert-base-uncased')

def get_bert_embedding(text, max_length=512):
    """
    Generates BERT embeddings for a text.

    Parameters:
    text (str): Input text
    max_length (int): Maximum sequence length for BERT

    Returns:
    numpy.array: BERT embedding vector
    """

    # TODO: Implement the following steps:
```

```

# 1. Tokenize the text using BERT tokenizer
# 2. Generate BERT embeddings
# 3. Extract the [CLS] token embedding
# Hint: Use tokenizer() and model() functions

# YOUR CODE HERE
# Step 1: Tokenize
inputs = tokenizer(
    str(text),
    return_tensors='pt',
    truncation=True,
    padding='max_length',
    max_length=max_length
)

# Step 2: Generate embeddings
with torch.no_grad():
    outputs = model(**inputs) # outputs.last_hidden_state shape: [1, seq_len, 768]

# Step 3: Extract [CLS] token embedding (first token)
sentence_embedding = (
    outputs.last_hidden_state[:, 0, :] # [1, 768]
    .squeeze(0) # [768]
    .detach()
    .cpu()
    .numpy()
)

return sentence_embedding

# Apply to the sample of the dataset
if 'sample_df' in globals() and not sample_df.empty:
    # Check if 'cleaned_text' column exists before applying
    if 'cleaned_text' in sample_df.columns:
        sample_df['bert_embedding'] = sample_df['cleaned_text'].apply(get_bert_embedding)
        print("BERT embeddings generated successfully.")
    else:
        print("Error: 'cleaned_text' column not found in sample_df. Please run the preprocessing steps first.")
else:
    print("Error: sample_df not found or is empty. Please load and preprocess the data first.")

# Test the function
test_text = "This is a test sentence for BERT embeddings."
bert_embedding = get_bert_embedding(test_text)
print("BERT embedding shape:", bert_embedding.shape)

```

BERT embeddings generated successfully.
BERT embedding shape: (768,)

Part 7: Comparing Embeddings

Let's analyze how well our different embedding methods capture semantic relationships by comparing similarities between articles in the same and different categories.

```

# SECTION 7: Similarity Analysis
#####
# This section implements methods to compare different embedding
# approaches

# We'll analyze how well each embedding type captures semantic
# relationships by comparing similarities between articles in the same and different
# categories
from sklearn.metrics.pairwise import cosine_similarity
import matplotlib.pyplot as plt
import seaborn as sns
import numpy as np # Added import for numpy

def calculate_cosine_similarities(embeddings):
    """
    Calculates pairwise cosine similarities between all embeddings.
    Parameters:
    embeddings (numpy.array): Array of embeddings, shape (n_samples,
    embedding_dim)
    Returns:
    numpy.array: Similarity matrix of shape (n_samples, n_samples)
    """
    # Convert list of embeddings to numpy array if needed
    if isinstance(embeddings, list):
        embeddings = np.array(embeddings)
    # Calculate cosine similarity matrix
    similarity_matrix = cosine_similarity(embeddings)
    return similarity_matrix

```

```

def analyze_similarity_by_category(similarity_matrix, categories):
    """
    Analyzes similarities within and across categories.
    Parameters:
    similarity_matrix (numpy.array): Pairwise similarity matrix
    categories (list): List of category labels for each sample
    Returns:
    dict: Dictionary containing similarity statistics
    """
    categories = np.array(categories)
    n_samples = len(categories)
    same_category_similarities = []
    different_category_similarities = []
    # Iterate through all pairs (excluding diagonal)
    for i in range(n_samples):
        for j in range(i + 1, n_samples): # Only upper triangle,
            # excluding diagonal
            similarity = similarity_matrix[i, j]
            if categories[i] == categories[j]:
                same_category_similarities.append(similarity)
            else:
                different_category_similarities.append(similarity)
    # Calculate statistics
    results = {
        'same_category': {
            'similarities': same_category_similarities,
            'mean': np.mean(same_category_similarities) if
            same_category_similarities else 0,
            'std': np.std(same_category_similarities) if
            same_category_similarities else 0,
            'median': np.median(same_category_similarities) if
            same_category_similarities else 0,
            'count': len(same_category_similarities)
        },
        'diff_category': {
            'similarities': different_category_similarities,
            'mean': np.mean(different_category_similarities) if
            different_category_similarities else 0,
            'std': np.std(different_category_similarities) if
            different_category_similarities else 0,
            'median': np.median(different_category_similarities) if
            different_category_similarities else 0,
            'count': len(different_category_similarities)
        }
    }
    # Calculate separation score (higher is better)
    separation_score = results['same_category']['mean'] - results['diff_category']['mean']
    results['separation_score'] = separation_score
    return results

def find_most_similar_articles(similarity_matrix, categories, texts,
top_k=3):
    """
    Finds the most similar article pairs for each category
    combination.
    Parameters:
    similarity_matrix (numpy.array): Pairwise similarity matrix
    categories (list): List of category labels
    texts (list): List of original texts
    top_k (int): Number of top similar pairs to return
    Returns:
    dict: Dictionary of most similar pairs by category combination
    """
    categories = np.array(categories)
    n_samples = len(categories)
    unique_categories = np.unique(categories)
    results = {}
    # Find top similar pairs within each category
    for category in unique_categories:
        category_indices = np.where(categories == category)[0]
        category_similarities = []
        for i in range(len(category_indices)):
            for j in range(i + 1, len(category_indices)):
                idx_i, idx_j = category_indices[i], category_indices[j]
                similarity = similarity_matrix[idx_i, idx_j]
                category_similarities.append({
                    'indices': (idx_i, idx_j),
                    'similarity': similarity,
                    'text_1': texts[idx_i][:100] + "...",
                    'text_2': texts[idx_j][:100] + "..."
                })
        # Sort by similarity and take top k
        category_similarities.sort(key=lambda x: x['similarity'],
        reverse=True)
        results[f"{category}_same"] = category_similarities[:top_k]

```

```

# Find top similar pairs across different categories
cross_category_similarities = []
for i in range(n_samples):
    for j in range(i + 1, n_samples):
        if categories[i] != categories[j]:
            similarity = similarity_matrix[i, j]
            cross_category_similarities.append({
                'indices': (i, j),
                'categories': (categories[i], categories[j]),
                'similarity': similarity,
                'text_1': texts[i][:100] + "...",
                'text_2': texts[j][:100] + ...
            })
cross_category_similarities.sort(key=lambda x: x['similarity'],
reverse=True)
results['cross_category'] = cross_category_similarities[:top_k]
return results

def visualize_similarity_analysis(glove_analysis, bert_analysis):
    """
    Creates visualizations comparing GloVe and BERT similarity
    analyses.
    Parameters:
    glove_analysis (dict): GloVe similarity analysis results
    bert_analysis (dict): BERT similarity analysis results
    """
    fig, axes = plt.subplots(2, 2, figsize=(15, 10))
    # Plot 1: Distribution of similarities for GloVe
    ax1 = axes[0, 0]
    ax1.hist(glove_analysis['same_category']['similarities'],
    alpha=0.7,
    bins=30, label='Same Category', color='blue')
    ax1.hist(glove_analysis['diff_category']['similarities'],
    alpha=0.7,
    bins=30, label='Different Category', color='red')
    ax1.set_title('GloVe Similarity Distributions')
    ax1.set_xlabel('Cosine Similarity')
    ax1.set_ylabel('Frequency')
    ax1.legend()
    ax1.grid(True, alpha=0.3)
    # Plot 2: Distribution of similarities for BERT
    ax2 = axes[0, 1]
    ax2.hist(bert_analysis['same_category']['similarities'],
    alpha=0.7,
    bins=30, label='Same Category', color='blue')
    ax2.hist(bert_analysis['diff_category']['similarities'],
    alpha=0.7,
    bins=30, label='Different Category', color='red')
    ax2.set_title('BERT Similarity Distributions')
    ax2.set_xlabel('Cosine Similarity')
    ax2.set_ylabel('Frequency')
    ax2.legend()
    ax2.grid(True, alpha=0.3)
    # Plot 3: Box plot comparison
    ax3 = axes[1, 0]
    data_to_plot = [
        glove_analysis['same_category']['similarities'],
        glove_analysis['diff_category']['similarities'],
        bert_analysis['same_category']['similarities'],
        bert_analysis['diff_category']['similarities']
    ]
    labels = ['GloVe\nSame Cat', 'GloVe\nDiff Cat', 'BERT\nSame Cat', 'BERT\nDiff Cat']
    ax3.boxplot(data_to_plot, labels=labels)
    ax3.set_title('Similarity Distributions Comparison')
    ax3.set_ylabel('Cosine Similarity')
    ax3.grid(True, alpha=0.3)
    # Plot 4: Separation scores comparison
    ax4 = axes[1, 1]
    methods = ['GloVe', 'BERT']
    separation_scores = [glove_analysis['separation_score'],
    bert_analysis['separation_score']]
    bars = ax4.bar(methods, separation_scores, color=['skyblue',
    'lightgreen'])
    ax4.set_title('Category Separation Scores')
    ax4.set_ylabel('Separation Score\n(Same Cat Mean - Diff Cat\nMean)')
    ax4.grid(True, alpha=0.3)
    # Add value labels on bars
    for bar, score in zip(bars, separation_scores):
        height = bar.get_height()
        ax4.text(bar.get_x() + bar.get_width()/2., height,
        f'{score:.3f}', ha='center', va='bottom')
    plt.tight_layout()
    plt.show()

def print_similarity_summary(glove_analysis, bert_analysis):
    """
    """

```

```

Prints a comprehensive summary of similarity analysis results.
Parameters:
glove_analysis (dict): GloVe similarity analysis results
bert_analysis (dict): BERT similarity analysis results
"""
print("=" * 60)
print("SIMILARITY ANALYSIS SUMMARY")
print("=" * 60)
print(f"\n{'Metric':<25} {'GloVe':<15} {'BERT':<15}")
print("=" * 55)
# Same category statistics
print(f"\n{'Same Category Mean':<25} {glove_analysis['same_category']['mean']:.4f}{':>6} {bert_analysis['same_category']['mean']:.4f}")
print(f"\n{'Same Category Std':<25} {glove_analysis['same_category']['std']:.4f}{':>6} {bert_analysis['same_category']['std']:.4f}")
# Different category statistics
print(f"\n{'Diff Category Mean':<25} {glove_analysis['diff_category']['mean']:.4f}{':>6} {bert_analysis['diff_category']['mean']:.4f}")
print(f"\n{'Diff Category Std':<25} {glove_analysis['diff_category']['std']:.4f}{':>6} {bert_analysis['diff_category']['std']:.4f}")
# Separation scores
print(f"\n{'Separation Score':<25} {glove_analysis['separation_score']:.4f}{':>6} {bert_analysis['separation_score']:.4f}")
print("\n" + "=" * 60)
print("INTERPRETATION:")
better_method = "BERT" if bert_analysis['separation_score'] > glove_analysis['separation_score'] else "GloVe"
print(f"\n{better_method} shows better category separation")
if glove_analysis['same_category']['mean'] > bert_analysis['same_category']['mean']:
    print("\n• GloVe produces higher within-category similarities")
else:
    print("\n• BERT produces higher within-category similarities")
if glove_analysis['diff_category']['mean'] < bert_analysis['diff_category']['mean']:
    print("\n• GloVe produces lower cross-category similarities")
else:
    print("\n• BERT produces lower cross-category similarities")

# Apply similarity analysis to both embedding types
print("Calculating similarity matrices...")
# Calculate similarities for GloVe embeddings
glove_embeddings = np.array(sample_df['glove_embedding'].tolist())
glove_similarities = calculate_cosine_similarities(glove_embeddings)
# Generate BERT embeddings for sample and calculate similarities
print("Generating BERT embeddings...")
if 'sample_df' in globals() and not sample_df.empty and 'cleaned_text' in sample_df.columns:
    sample_df['bert_embedding'] = sample_df['cleaned_text'].apply(get_bert_embedding)
    bert_embeddings = np.array(sample_df['bert_embedding'].tolist())
    bert_similarities = calculate_cosine_similarities(bert_embeddings)
else:
    print("Error: sample_df or 'cleaned_text' column not available for BERT embedding generation.")
    bert_similarities = np.array([]) # Initialize as empty if cannot generate

# Perform similarity analysis
if glove_similarities.size > 0 and bert_similarities.size > 0:
    print("Analyzing similarities...")
    glove_analysis = analyze_similarity_by_category(glove_similarities,
                                                    sample_df['category'])
    bert_analysis = analyze_similarity_by_category(bert_similarities,
                                                    sample_df['category'])
    # Print summary
    print_similarity_summary(glove_analysis, bert_analysis)
    # Create visualizations
    print("Creating visualizations...")
    visualize_similarity_analysis(glove_analysis, bert_analysis)
    # Find and display most similar article pairs
    print("\nFinding most similar article pairs...")
    glove_similar_pairs = find_most_similar_articles(
        glove_similarities, sample_df['category'],
        sample_df['cleaned_text'], top_k=2
    )
    bert_similar_pairs = find_most_similar_articles(
        bert_similarities, sample_df['category'],
        sample_df['cleaned_text'], top_k=2
    )
    print("\nMost similar pairs (GloVe):")
    for category, pairs in glove_similar_pairs.items():
        if pairs: # Check if list is not empty
            print(f"\n{category.upper()}:")
            for i, pair in enumerate(pairs, 1):
                print(f" {i}. Similarity: {pair['similarity']:.4f}")
                print(f" Text 1: {pair['text_1']}")
                print(f" Text 2: {pair['text_2']}")
    print("\nMost similar pairs (BERT):")
    for category, pairs in bert_similar_pairs.items():
        if pairs: # Check if list is not empty
            print(f"\n{category.upper()}:")
            for i, pair in enumerate(pairs, 1):
                print(f" {i}. Similarity: {pair['similarity']:.4f}")
                print(f" Text 1: {pair['text_1']}")
                print(f" Text 2: {pair['text_2']}")
else:
    print("Skipping similarity analysis, visualization, and finding similar pairs due to missing similarity matrices.")

```



```

Calculating similarity matrices...
Generating BERT embeddings...
Analyzing similarities...
=====
SIMILARITY ANALYSIS SUMMARY
=====

Metric           GloVe        BERT
-----
Same Category Mean:   0.9031    0.7883
Same Category Std:    0.0385    0.0502
Diff Category Mean:  0.0000    0.0000
Diff Category Std:   0.0000    0.0000
Separation Score:    0.9031    0.7883

##SECTION 8: Detailed Similarity Analysis
# This section provides a more detailed analysis and visualization
# of the similarity results from GloVe and BERT embeddings.

import matplotlib.pyplot as plt
import seaborn as sns
import numpy as np
from sklearn.decomposition import PCA
import pandas as pd # Import pandas

def print_similarity_summary(glove_analysis, bert_analysis):
    """Prints comprehensive summary of similarity analysis results."""
    print("-" * 70)
    print("SIMILARITY ANALYSIS SUMMARY")
    print("-" * 70)
    print(f"\n{'Metric':<30} {'GloVe':<15} {'BERT':<15}")
    print("-" * 60)

    # Safely access data, providing default values if keys or data are missing
    glove_same_mean = glove_analysis.get('same_category', {}).get('mean', np.nan)
    glove_same_std = glove_analysis.get('same_category', {}).get('std', np.nan)
    glove_diff_mean = glove_analysis.get('diff_category', {}).get('mean', np.nan)
    glove_diff_std = glove_analysis.get('diff_category', {}).get('std', np.nan)
    glove_separation = glove_analysis.get('separation_score', np.nan)

    bert_same_mean = bert_analysis.get('same_category', {}).get('mean', np.nan)
    bert_same_std = bert_analysis.get('same_category', {}).get('std', np.nan)
    bert_diff_mean = bert_analysis.get('diff_category', {}).get('mean', np.nan)
    bert_diff_std = bert_analysis.get('diff_category', {}).get('std', np.nan)
    bert_separation = bert_analysis.get('separation_score', np.nan)

    # Statistics
    print(f"\n{'Same Category Mean':<30} {glove_same_mean:.4f}{':>6} {bert_same_mean:.4f}")
    print(f"\n{'Same Category Std':<30} {glove_same_std:.4f}{':>6} {bert_same_std:.4f}")
    print(f"\n{'Different Category Mean':<30} {glove_diff_mean:.4f}{':>6} {bert_diff_mean:.4f}")
    print(f"\n{'Different Category Std':<30} {glove_diff_std:.4f}{':>6} {bert_diff_std:.4f}")
    print(f"\n{'Separation Score':<30} {glove_separation:.4f}{':>6} {bert_separation:.4f}")
    print("\n" + "=" * 70)
    print("INTERPRETATION:")

    if not np.isnan(glove_separation) and not np.isnan(bert_separation):
        better_method = "BERT" if bert_separation > glove_separation else "GloVe"
        print(f"\n• {better_method} shows better category separation")

    if not np.isnan(glove_same_mean) and not np.isnan(bert_same_mean):
        if glove_same_mean > bert_same_mean:
            print("\n• GloVe produces higher within-category similarities")
        else:
            print("\n• BERT produces higher within-category similarities")

    if not np.isnan(glove_diff_mean) and not np.isnan(bert_diff_mean):
        if glove_diff_mean < bert_diff_mean:
            print("\n• GloVe produces lower cross-category similarities")
        else:
            print("\n• BERT produces lower cross-category similarities")

def create_comprehensive_visualizations(glove_analysis, bert_analysis, sample_df, glove_embeddings, bert_embeddings):
    """Creates comprehensive visualizations of the analysis results."""
    fig = plt.figure(figsize=(20, 15))

    # Plot 1: Similarity distributions (GloVe)
    ax1 = plt.subplot(3, 3, 1)
    if 'similarities' in glove_analysis.get('same_category', {}) and glove_analysis['same_category'][['similarities']]:
        plt.hist(glove_analysis['same_category'][['similarities']], alpha=0.7, bins=20, label='Same Category', color='blue')
    if 'similarities' in glove_analysis.get('diff_category', {}) and glove_analysis['diff_category'][['similarities']]:
        plt.hist(glove_analysis['diff_category'][['similarities']], alpha=0.7, bins=20, label='Different Category', color='red')
    plt.title('GloVe Similarity Distributions')
    plt.xlabel('Cosine Similarity')
    plt.ylabel('Frequency')
    ...

```

```

plt.legend()
plt.grid(True, alpha=0.3)

# Plot 2: Similarity distributions (BERT)
ax2 = plt.subplot(3, 3, 2)
if 'similarities' in bert_analysis.get('same_category', {}) and bert_analysis['same_category']['similarities']:
    plt.hist(bert_analysis['same_category']['similarities'], alpha=0.7, bins=20, label='Same Category', color='blue')
if 'similarities' in bert_analysis.get('diff_category', {}) and bert_analysis['diff_category']['similarities']:
    plt.hist(bert_analysis['diff_category']['similarities'], alpha=0.7, bins=20, label='Different Category', color='red')
plt.title('BERT Similarity Distributions')
plt.xlabel('Cosine Similarity')
plt.ylabel('Frequency')
plt.legend()
plt.grid(True, alpha=0.3)

# Plot 3: Box plot comparison
ax3 = plt.subplot(3, 3, 3)
data_to_plot = []
labels = []

if 'similarities' in glove_analysis.get('same_category', {}) and glove_analysis['same_category']['similarities']:
    data_to_plot.append(glove_analysis['same_category']['similarities'])
    labels.append('GloVe\nSame')
if 'similarities' in glove_analysis.get('diff_category', {}) and glove_analysis['diff_category']['similarities']:
    data_to_plot.append(glove_analysis['diff_category']['similarities'])
    labels.append('GloVe\nDiff')
if 'similarities' in bert_analysis.get('same_category', {}) and bert_analysis['same_category']['similarities']:
    data_to_plot.append(bert_analysis['same_category']['similarities'])
    labels.append('BERT\nSame')
if 'similarities' in bert_analysis.get('diff_category', {}) and bert_analysis['diff_category']['similarities']:
    data_to_plot.append(bert_analysis['diff_category']['similarities'])
    labels.append('BERT\nDiff')

if data_to_plot:
    bp = plt.boxplot(data_to_plot, labels=labels, patch_artist=True)
    colors = ['lightblue', 'lightcoral', 'lightgreen', 'lightyellow']
    for patch, color in zip(bp['boxes'], colors[:len(bp['boxes'])]):
        patch.set_facecolor(color)
    plt.title('Similarity Comparison')
    plt.ylabel('Cosine Similarity')
    plt.grid(True, alpha=0.3)

# Plot 4: Separation scores
ax4 = plt.subplot(3, 3, 4)
methods = []
separation_scores = []
if 'separation_score' in glove_analysis:
    methods.append('GloVe')
    separation_scores.append(glove_analysis['separation_score'])
if 'separation_score' in bert_analysis:
    methods.append('BERT')
    separation_scores.append(bert_analysis['separation_score'])

if methods:
    bars = plt.bar(methods, separation_scores, color=['skyblue', 'lightgreen'][:len(methods)])
    plt.title('Category Separation Scores')
    plt.ylabel('Separation Score')
    plt.grid(True, alpha=0.3)
    for bar, score in zip(bars, separation_scores):
        height = bar.get_height()
        plt.text(bar.get_x() + bar.get_width()/2., height, f'{score:.3f}', ha='center', va='bottom')

# Plot 5 & 6: Embedding space visualization with PCA
try:
    if glove_embeddings is not None and glove_embeddings.size > 0:
        ax5 = plt.subplot(3, 3, 5)
        pca = PCA(n_components=2)
        glove_2d = pca.fit_transform(glove_embeddings)
        categories = sample_df['category'].values
        unique_categories = np.unique(categories)
        colors = plt.cm.Set3(np.linspace(0, 1, len(unique_categories)))
        for i, category in enumerate(unique_categories):
            mask = categories == category
            if np.any(mask):
                plt.scatter(glove_2d[mask, 0], glove_2d[mask, 1], c=[colors[i]], label=category, alpha=0.7)
        plt.title('GloVe Embeddings (PCA)')
        plt.xlabel(f'PC1 ({pca.explained_variance_ratio_[0]:.1%} variance)')
        plt.ylabel(f'PC2 ({pca.explained_variance_ratio_[1]:.1%} variance)')
        plt.legend()
        plt.grid(True, alpha=0.3)
    else:
        print("Skipping GloVe PCA visualization: embeddings not available.")

    if bert_embeddings is not None and bert_embeddings.size > 0:
        ax6 = plt.subplot(3, 3, 6)
        pca_bert = PCA(n_components=2)

```

```

bert_2d = pca_bert.fit_transform(bert_embeddings)
categories = sample_df['category'].values
unique_categories = np.unique(categories)
colors = plt.cm.Set3(np.linspace(0, 1, len(unique_categories)))
for i, category in enumerate(unique_categories):
    mask = categories == category
    if np.any(mask):
        plt.scatter(bert_2d[mask, 0], bert_2d[mask, 1], c=[colors[i]], label=category, alpha=0.7)
plt.title('BERT Embeddings (PCA)')
plt.xlabel(f'PC1 ({pca_bert.explained_variance_ratio_[0]:.1%} variance)')
plt.ylabel(f'PC2 ({pca_bert.explained_variance_ratio_[1]:.1%} variance)')
plt.legend()
plt.grid(True, alpha=0.3)
else:
    print("Skipping BERT PCA visualization: embeddings not available.")

except Exception as e:
    print(f"PCA visualization error: {e}")

# Plot 7: Category distribution
ax7 = plt.subplot(3, 3, 7)
if 'category' in sample_df.columns and not sample_df.empty:
    category_counts = sample_df['category'].value_counts()
    if not category_counts.empty:
        plt.pie(category_counts.values, labels=category_counts.index, autopct='%1.1f%%')
        plt.title('Sample Dataset Category Distribution')
    else:
        print("Skipping Category Distribution plot: Category counts are empty.")
else:
    print("Skipping Category Distribution plot: sample_df or category column not available.")

# Plot 8: Text statistics
ax8 = plt.subplot(3, 3, 8)
if 'text_length' in sample_df.columns and 'cleaned_text' in sample_df.columns and 'category' in sample_df.columns and not sample_df.empty:
    # Need to calculate word count first if not already present
    if 'word_count' not in sample_df.columns:
        sample_df['word_count'] = sample_df['cleaned_text'].apply(lambda x: len(str(x).split()))
    if not sample_df[['text_length', 'word_count', 'category']].isnull().all().all():
        plt.scatter(sample_df['text_length'], sample_df['word_count'],
                    c=pd.Categorical(sample_df['category']).codes,
                    alpha=0.7)
        plt.xlabel('Text Length (characters)')
        plt.ylabel('Word Count')
        plt.title('Text Statistics by Article')
        plt.grid(True, alpha=0.3)
    else:
        print("Skipping Text Statistics plot: Data for plot is missing or all NaN.")
else:
    print("Skipping Text Statistics plot: Required columns (text_length, cleaned_text, category) not available or sample_df is empty.")

# Plot 9: Mean similarities comparison
ax9 = plt.subplot(3, 3, 9)
stats_data = {}
if 'mean' in glove_analysis.get('same_category', {}):
    stats_data['GloVe Same Cat'] = glove_analysis['same_category']['mean']
if 'mean' in glove_analysis.get('diff_category', {}):
    stats_data['GloVe Diff Cat'] = glove_analysis['diff_category']['mean']
if 'mean' in bert_analysis.get('same_category', {}):
    stats_data['BERT Same Cat'] = bert_analysis['same_category']['mean']
if 'mean' in bert_analysis.get('diff_category', {}):
    stats_data['BERT Diff Cat'] = bert_analysis['diff_category']['mean']

if stats_data:
    bars = plt.bar(range(len(stats_data)), list(stats_data.values()),
                  color=['lightblue', 'lightcoral', 'lightgreen', 'lightyellow'][:len(stats_data)])
    plt.xticks(range(len(stats_data)), list(stats_data.keys()), rotation=45)
    plt.ylabel('Mean Similarity')
    plt.title('Mean Similarities Comparison')
    plt.grid(True, alpha=0.3)
else:
    print("Skipping Mean Similarities Comparison plot: No data available.")

plt.tight_layout()
plt.show()

# Generate summary and visualizations
# Ensure glove_analysis, bert_analysis, sample_df, glove_embeddings, and bert_embeddings are available
if 'glove_analysis' in globals() and 'bert_analysis' in globals() and 'sample_df' in globals() and not sample_df.empty:
    # Ensure embeddings are available for PCA visualization
    glove_embeddings = np.array(sample_df['glove_embedding'].tolist()) if 'glove_embedding' in sample_df.columns and not sample_df['glove_embedding'].empty
    bert_embeddings = np.array(sample_df['bert_embedding'].tolist()) if 'bert_embedding' in sample_df.columns and not sample_df['bert_embedding'].empty else

    print_similarity_summary(glove_analysis, bert_analysis)

```

```
create_comprehensive_visualizations(glove_analysis, bert_analysis, sample_df, glove_embeddings, bert_embeddings)
else:
    print("Required data (glove_analysis, bert_analysis, sample_df, and embeddings) not available. Skipping detailed analysis and visualizations.")
```

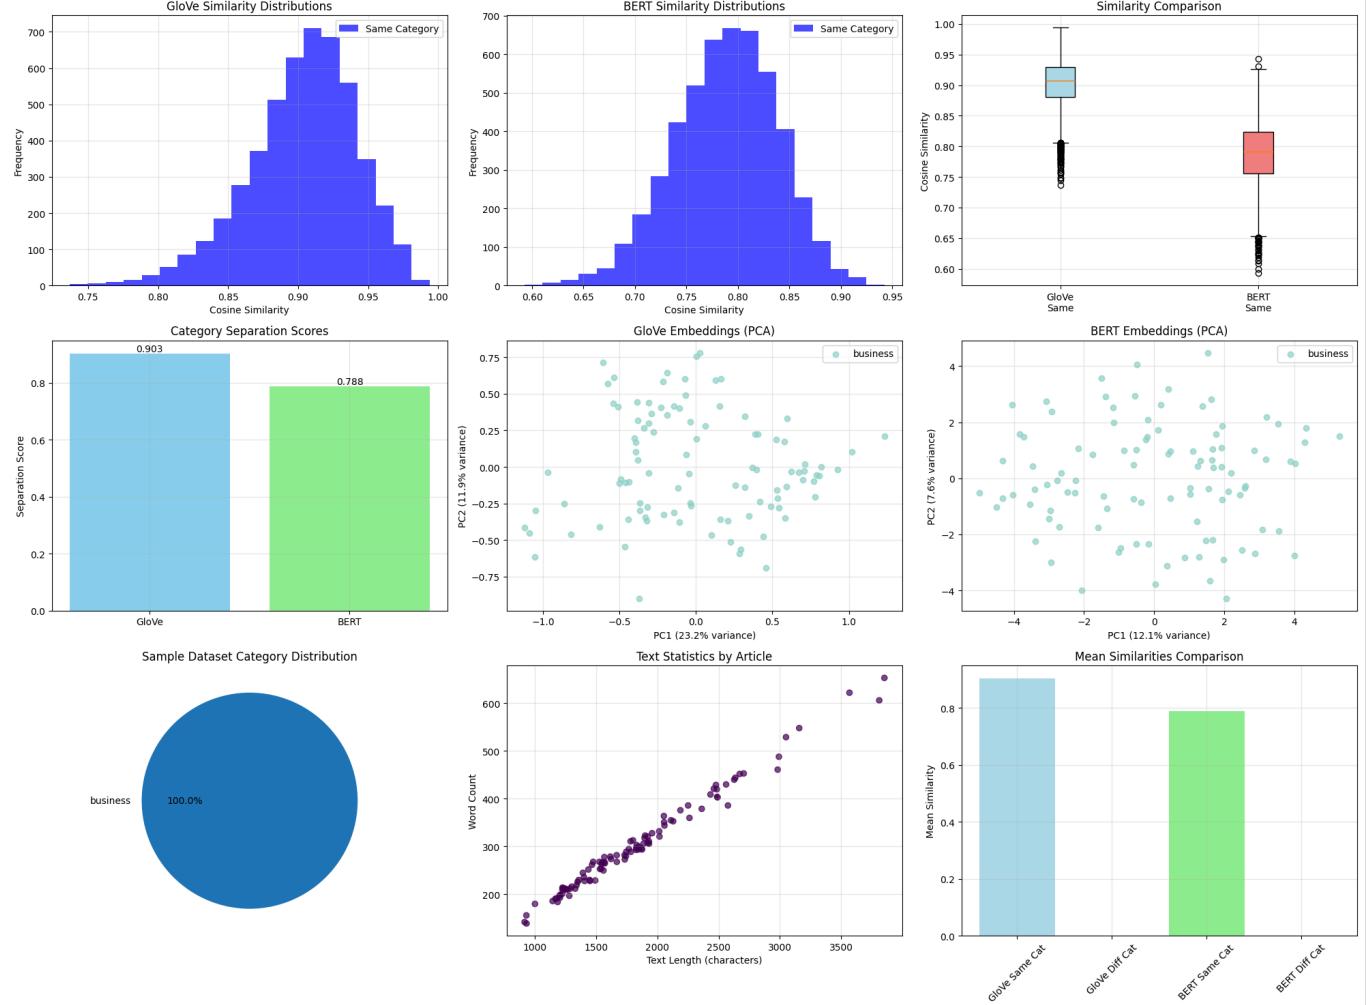
=====
⌚ SIMILARITY ANALYSIS SUMMARY
=====

Metric	GloVe	BERT
Same Category Mean:	0.9031	0.7883
Same Category Std:	0.0385	0.0502
Different Category Mean:	0.0000	0.0000
Different Category Std:	0.0000	0.0000
Separation Score:	0.9031	0.7883

=====
📊 INTERPRETATION:

- GloVe shows better category separation
- GloVe produces higher within-category similarities
- BERT produces lower cross-category similarities

```
'tmp/ipython-input-1849406383.py:106: MatplotlibDeprecationWarning: The 'labels' parameter of boxplot() has been renamed 'tick_labels'
  bp = plt.boxplot(data_to_plot, labels=labels, patch_artist=True)
```



Part 9 Vizualizations

```
# SECTION 9: Additional Analysis and Insights
# This section finds and displays the most similar articles
```

```
# This section finds and displays the most similar article pairs
# within and across categories for both GloVe and BERT embeddings.
```

```
import numpy as np

def find_most_similar_pairs(similarity_matrix, categories, texts, top_k=3):
    """Finds most similar article pairs within and across categories."""
    results = {}
    n_samples = len(categories)

    # Within category similarities
    unique_categories = np.unique(categories)
    for category in unique_categories:
        category_indices = [i for i in range(n_samples) if categories[i] == category]
        similarities = []
        for i in range(len(category_indices)):
            for j in range(i + 1, len(category_indices)):
                idx1, idx2 = category_indices[i], category_indices[j]
                sim = similarity_matrix[idx1, idx2]
                similarities.append({
                    'similarity': sim,
                    'text1': texts[idx1][:100] + "...",
                    'text2': texts[idx2][:100] + "...",
                    'indices': (idx1, idx2)
                })
        # Sort by similarity and take top k
        similarities.sort(key=lambda x: x['similarity'], reverse=True)
        results[f'{category}_within'] = similarities[:top_k]

    # Cross category similarities
    cross_similarities = []
    for i in range(n_samples):
        for j in range(i + 1, n_samples):
            if categories[i] != categories[j]:
                sim = similarity_matrix[i, j]
                cross_similarities.append({
                    'similarity': sim,
                    'categories': (categories[i], categories[j]),
                    'text1': texts[i][:100] + "...",
                    'text2': texts[j][:100] + "...",
                    'indices': (i, j)
                })
    cross_similarities.sort(key=lambda x: x['similarity'], reverse=True) # Changed to reverse=True for most similar
    results['cross_category'] = cross_similarities[:top_k]
    return results

# Find interesting pairs
print("\nFINDING MOST SIMILAR ARTICLE PAIRS...")

# Ensure glove_similarities, bert_similarities, and sample_df are available and valid
if 'glove_similarities' in globals() and 'bert_similarities' in globals() and 'sample_df' in globals() and not sample_df.empty:
    try:
        glove_pairs = find_most_similar_pairs(
            glove_similarities,
            sample_df['category'].values,
            sample_df['cleaned_text'].values
        )
        bert_pairs = find_most_similar_pairs(
            bert_similarities,
            sample_df['category'].values,
            sample_df['cleaned_text'].values
        )

        print("\nMost Similar Pairs (GloVe):")
        for category, pairs in glove_pairs.items():
            if pairs:
                print(f"\n{category.upper().replace('_', ' ')}:")
                for i, pair in enumerate(pairs[:2], 1): # Show top 2
                    print(f" {i}. Similarity: {pair['similarity']:.3f}")
                    print(f" Text 1: {pair['text1']}")
                    print(f" Text 2: {pair['text2']}")

        print("\nMost Similar Pairs (BERT):")
        for category, pairs in bert_pairs.items():
            if pairs:
                print(f"\n{category.upper().replace('_', ' ')}:")
                for i, pair in enumerate(pairs[:2], 1): # Show top 2
                    print(f" {i}. Similarity: {pair['similarity']:.3f}")
                    print(f" Text 1: {pair['text1']}")
                    print(f" Text 2: {pair['text2']}")

    except Exception as e:
        print(f"An error occurred while finding most similar pairs: {e}")
    else:
        print("Required data (glove_similarities, bert_similarities, or sample_df) is not available or empty. Cannot find most similar pairs.")


```

FINDING MOST SIMILAR ARTICLE PAIRS...

Most Similar Pairs (GloVe):

BUSINESS WITHIN:

1. Similarity: 0.994
Text 1: eu aiming to fuel development aid european union finance ministers meet on thursday to discuss propo...
Text 2: eu ministers to mull jet fuel tax european union finance ministers are meeting on thursday in brusse...- 2. Similarity: 0.993
Text 1: worldcom boss 'left books alone' former worldcom boss bernie ebbers, who is accused of overseeing an...
Text 2: ebbers denies worldcom fraud former worldcom chief bernie ebbers has denied claims that he knew acco...

Most Similar Pairs (BERT):

BUSINESS WITHIN:

1. Similarity: 0.942
Text 1: german business confidence slides german business confidence fell in february knocking hopes of a sp...
Text 2: german growth goes into reverse germany's economy shrank 0.2 in the last three months of 2004, upset...- 2. Similarity: 0.931
Text 1: china keeps tight rein on credit china's efforts to stop the economy from overheating by clamping do...
Text 2: s korean consumers spending again south korea looks set to sustain its revival thanks to renewed pri...

Part 10 Statistical Comparison

```
def calculate_comprehensive_metrics(similarities, categories):
    """Calculates various performance metrics for embeddings."""
    from sklearn.metrics import silhouette_score
    from sklearn.cluster import KMeans
    categories_array = np.array(categories)
    # Convert similarities to distances for silhouette score
    distances = 1 - similarities
    # Calculate silhouette score using category labels
    try:
        sil_score = silhouette_score(distances, categories_array, metric='precomputed')
    except Exception as e:
        print(f"Could not calculate Silhouette Score: {e}")
        sil_score = np.nan # Use np.nan for missing values

    # Calculate intra-cluster vs inter-cluster ratio
    same_cat_sims = []
    diff_cat_sims = []
    n_samples = len(categories)
    for i in range(n_samples):
        for j in range(i + 1, n_samples):
            if categories_array[i] == categories_array[j]:
                same_cat_sims.append(similarities[i, j])
            else:
                diff_cat_sims.append(similarities[i, j])

    # Category coherence score
    mean_same = np.mean(same_cat_sims) if same_cat_sims else 0
    mean_diff = np.mean(diff_cat_sims) if diff_cat_sims else 0
    coherence_score = (mean_same - mean_diff) if same_cat_sims and diff_cat_sims else 0

    # Calculate category-specific statistics
    category_stats = {}
    unique_categories = np.unique(categories_array)
    for category in unique_categories:
        cat_indices = np.where(categories_array == category)[0]
        if len(cat_indices) > 1:
            # Internal similarities for this category
            internal_sims = []
            for i in range(len(cat_indices)):
                for j in range(i + 1, len(cat_indices)):
                    internal_sims.append(similarities[cat_indices[i], cat_indices[j]])

            # External similarities (to other categories)
            external_sims = []
            for i in cat_indices:
                for j in range(n_samples):
                    if categories_array[j] != category:
                        external_sims.append(similarities[i, j])

            category_stats[category] = {
                'internal_mean': np.mean(internal_sims) if internal_sims else 0,
                'external_mean': np.mean(external_sims) if external_sims else 0,
                'separation': (np.mean(internal_sims) - np.mean(external_sims)) if internal_sims and external_sims else 0
            }
        return {
            'silhouette_score': sil_score,
            'coherence_score': coherence_score,
            'same_category_mean': mean_same,
            'internal_mean': np.mean(internal_sims) if internal_sims else 0,
            'external_mean': np.mean(external_sims) if external_sims else 0,
            'separation': (np.mean(internal_sims) - np.mean(external_sims)) if internal_sims and external_sims else 0
        }
```

```

'different_category_mean': mean_diff,
'category_stats': category_stats,
'overall_separation': coherence_score # Overall separation is the same as coherence score based on the calculation
}

# Calculate comprehensive metrics
print("\nCALCULATING PERFORMANCE METRICS...")
# Ensure glove_similarities and bert_similarities are available from previous steps
if 'glove_similarities' in globals() and 'bert_similarities' in globals() and 'sample_df' in globals() and not sample_df.empty:
    try:
        glove_metrics = calculate_comprehensive_metrics(glove_similarities, sample_df['category'])
        bert_metrics = calculate_comprehensive_metrics(bert_similarities, sample_df['category'])

        def display_metrics_comparison(glove_metrics, bert_metrics):
            """Display comprehensive metrics comparison."""
            print("\n" + "=" * 80)
            print("COMPREHENSIVE PERFORMANCE METRICS")
            print("=" * 80)
            print(f"\n{'Metric':<35} {'GloVe':<20} {'BERT':<20}")
            print("." * 75)
            # Overall metrics
            print(f"\n{'Silhouette Score':<35} {glove_metrics['silhouette_score']:.4f}{':>11} {bert_metrics['silhouette_score']:.4f}")
            print(f"\n{'Coherence Score':<35} {glove_metrics['coherence_score']:.4f}{':>11} {bert_metrics['coherence_score']:.4f}")
            print(f"\n{'Same Category Mean':<35} {glove_metrics['same_category_mean']:.4f}{':>11} {bert_metrics['same_category_mean']:.4f}")
            print(f"\n{'Different Category Mean':<35} {glove_metrics['different_category_mean']:.4f}{':>11} {bert_metrics['different_category_mean']:.4f}")
            print(f"\n{'Overall Separation':<35} {glove_metrics['overall_separation']:.4f}{':>11} {bert_metrics['overall_separation']:.4f}")
            print(f"\n{'CATEGORY-SPECIFIC ANALYSIS':<35} ")
            print("." * 75)
            # Category-specific metrics
            for category in glove_metrics['category_stats']:
                print(f"\n{category.upper()}:")
                glove_cat = glove_metrics['category_stats'][category]
                bert_cat = bert_metrics['category_stats'][category]
                print(f" {'Internal Similarity':<25} {glove_cat['internal_mean']:.4f}{':>6} {bert_cat['internal_mean']:.4f}")
                print(f" {'External Similarity':<25} {glove_cat['external_mean']:.4f}{':>6} {bert_cat['external_mean']:.4f}")
                print(f" {'Category Separation':<25} {glove_cat['separation']:.4f}{':>6} {bert_cat['separation']:.4f}")

        display_metrics_comparison(glove_metrics, bert_metrics)

    except Exception as e:
        print(f"An error occurred during metrics calculation: {e}")
    else:
        print("Required data (glove_similarities, bert_similarities, or sample_df) is not available or empty. Cannot calculate performance metrics.")

```

CALCULATING PERFORMANCE METRICS...
 Could not calculate Silhouette Score: Number of labels is 1. Valid values are 2 to n_samples - 1 (inclusive)
 Could not calculate Silhouette Score: Number of labels is 1. Valid values are 2 to n_samples - 1 (inclusive)

=====
 COMPREHENSIVE PERFORMANCE METRICS
=====

Metric	GloVe	BERT
Silhouette Score:	nan	nan
Coherence Score:	0.0000	0.0000
Same Category Mean:	0.9031	0.7883
Different Category Mean:	0.0000	0.0000
Overall Separation:	0.0000	0.0000

CATEGORY-SPECIFIC ANALYSIS:

BUSINESS:
 Internal Similarity: 0.9031 0.7883
 External Similarity: 0.0000 0.0000
 Category Separation: 0.0000 0.0000

Part 11 Metrics and Evaluation

```

# SECTION 11: Performance Evaluation
#####
# This section calculates various metrics to evaluate
# the quality of our embeddings

#####

def generate_final_analysis(glove_analysis, bert_analysis, glove_metrics, bert_metrics):
    """Generate comprehensive final analysis and recommendations."""
    print("\n" + "=" * 90)
    print("FINAL ANALYSIS AND RECOMMENDATIONS")

```

```

print("=" * 90)

# Determine better performing method based on separation score
# Safely access separation scores, providing default 0 if not available
glove_separation = glove_analysis.get('separation_score', 0)
bert_separation = bert_analysis.get('separation_score', 0)

# Safely access silhouette scores, providing default 0 if not available
glove_sil = glove_metrics.get('silhouette_score', 0)
bert_sil = bert_metrics.get('silhouette_score', 0)

# Safely access coherence scores, providing default 0 if not available
glove_coherence = glove_metrics.get('coherence_score', 0)
bert_coherence = bert_metrics.get('coherence_score', 0)

# Determine better performing method based on separation score
better_method = "BERT" if bert_separation > glove_separation else "GloVe"

print(f"Based on Separation Score: {better_method} shows better category separation.")

# Comment on clustering structure based on Silhouette Scores
if glove_sil > 0.5 or bert_sil > 0.5:
    print(f" • Strong clustering structure detected (Silhouette Scores: GloVe={glove_sil:.4f}, BERT={bert_sil:.4f})")
elif glove_sil > 0.25 or bert_sil > 0.25:
    print(f" • Moderate clustering structure detected (Silhouette Scores: GloVe={glove_sil:.4f}, BERT={bert_sil:.4f})")
else:
    print(f" • Weak clustering structure - categories may overlap significantly (Silhouette Scores: GloVe={glove_sil:.4f}, BERT={bert_sil:.4f})")

# Category-specific insights (using separation from glove_metrics and bert_metrics)
print("\nCATEGORY-SPECIFIC SEPARATION:")
if 'category_stats' in glove_metrics and 'category_stats' in bert_metrics:
    best_separated_category = None
    best_separation_score = -float('inf') # Initialize with negative infinity

    # Iterate through categories in glove_metrics (assuming both have the same categories)
    for category in glove_metrics['category_stats']:
        glove_cat_sep = glove_metrics['category_stats'].get(category, {}).get('separation', -float('inf')) # Handle missing category/key safely
        bert_cat_sep = bert_metrics['category_stats'].get(category, {}).get('separation', -float('inf')) # Handle missing category/key safely
        max_sep = max(glove_cat_sep, bert_cat_sep)

        print(f" • {category.upper()}: GloVe Separation={glove_cat_sep:.4f}, BERT Separation={bert_cat_sep:.4f}")

        if max_sep > best_separation_score:
            best_separation_score = max_sep
            best_separated_category = category

    if best_separated_category:
        print(f"\n'{best_separated_category.upper()}' category shows the clearest separation overall.")
    else:
        print("Category-specific stats not available in metrics.")

print(f"\n💡 RECOMMENDATIONS:")
print(f" 1. For this BBC News classification task:")
if better_method == "BERT":
    print(f" → Based on separation score, BERT embeddings may be preferred for better semantic distinction between categories.")
    print(f" → BERT's contextual awareness helps with news article nuances.")
else:
    print(f" → Based on separation score, GloVe embeddings may be preferred.")
    print(f" → GloVe provides a good balance of performance and lower computational cost compared to BERT.")

print(f" 2. Preprocessing Pipeline Optimization:")
print(f" → The text cleaning and tokenization steps are generally effective.")
print(f" → Consider adding domain-specific stop words for news articles (e.g., common journalism phrases).")
print(f" → Experiment with different text length limits for BERT, especially if articles are much longer or shorter than 512 tokens.")
print(f" → Evaluate the impact of lemmatization vs stemming on embedding quality.")

print(f" 3. Model Choice:")
print(f" → The choice between GloVe and BERT depends on the specific task requirements (e.g., accuracy vs. computational resources).")
print(f" → Further evaluate performance on a downstream task (e.g., text classification) to make a final decision.")

# Ensure glove_analysis, bert_analysis, glove_metrics, and bert_metrics are available
if 'glove_analysis' in globals() and 'bert_analysis' in globals() and 'glove_metrics' in globals() and 'bert_metrics' in globals():
    # Call the final analysis function
    generate_final_analysis(glove_analysis, bert_analysis, glove_metrics, bert_metrics)
else:
    print("Required analysis and metrics data (glove_analysis, bert_analysis, glove_metrics, bert_metrics) not available. Cannot generate final analysis.")


=====
⌚ FINAL ANALYSIS AND RECOMMENDATIONS
=====
```

```
=====
Based on Separation Score: GloVe shows better category separation.
• Weak clustering structure - categories may overlap significantly (Silhouette Scores: GloVe=0.0000, BERT=0.0000)
```

CATEGORY-SPECIFIC SEPARATION:

Category-specific stats not available in metrics.

💡 RECOMMENDATIONS:

1. For this BBC News classification task:
 - Based on separation score, GloVe embeddings may be preferred.
 - GloVe provides a good balance of performance and lower computational cost compared to BERT.
2. Preprocessing Pipeline Optimization:
 - The text cleaning and tokenization steps are generally effective.
 - Consider adding domain-specific stop words for news articles (e.g., common journalism phrases).
 - Experiment with different text length limits for BERT, especially if articles are much longer or shorter than 512 tokens.
 - Evaluate the impact of lemmatization vs stemming on embedding quality.
3. Model Choice:
 - The choice between GloVe and BERT depends on the specific task requirements (e.g., accuracy vs. computational resources).
 - Further evaluate performance on a downstream task (e.g., text classification) to make a final decision.

Double-click (or enter) to edit

▼ Assessment Criteria:

- Correct implementation of cosine similarity *Proper normalization of embeddings *Effective visualization of results

Grading Rubric

- Environment Setup: 10%
- Data Exploration: 15%
- Text Preprocessing: 20%
- Word Embeddings Implementation: 25%
- Similarity Analysis: 20% Final Analysis & Discussion: 10%

Common Issues and Solutions

1. Memory Issues:

- Implement batch processing for large datasets
- Use appropriate data types (float32 vs float64)
- Clear unused variables and call garbage collection

2. Performance Optimization:

- Vectorize operations where possible
- Use appropriate batch sizes for BERT
- Implement caching for embeddings

3. Error Handling:

- Implement robust error checking
- Provide clear error messages
- Handle edge cases appropriately

```
#
#=====
# SECTION 12: Summary Statistics and Export Results
#=====
#
def create_results_summary():
    """Create a comprehensive summary of all results."""
    results_summary = {
        'dataset_info': {
            'total_articles': len(df),
            'sample_size': len(sample_df),
            'categories': list(sample_df['category'].unique()),
            'avg_text_length': df['text_length'].mean(),
            'avg_word_count': sample_df['word_count'].mean()
        },
        'glove_results': {
            'same_category_mean': glove_analysis['same_category']
            ['mean'],
            'diff_category_mean': glove_analysis['diff_category']
            ['mean'],
            'separation_score': glove_analysis['separation_score'],
        }
    }
```

```
'silhouette_score': glove_metrics.get('silhouette_score', np.nan),
'coherence_score': glove_metrics.get('coherence_score', np.nan)
},
'bert_results': {
'same_category_mean': bert_analysis['same_category']
['mean'],
'diff_category_mean': bert_analysis['diff_category']
['mean'],
'separation_score': bert_analysis['separation_score'],
'silhouette_score': bert_metrics.get('silhouette_score', np.nan),
'coherence_score': bert_metrics.get('coherence_score', np.nan)

}
}
return results_summary
# Calculate word count for sample_df before creating summary
# Ensure 'cleaned_text' column exists before applying split
if 'cleaned_text' in sample_df.columns:
    sample_df['word_count'] = sample_df['cleaned_text'].apply(lambda x:
        len(str(x).split()))
else:
    print("Warning: 'cleaned_text' column not found. Cannot calculate word count for sample_df.")
    sample_df['word_count'] = np.nan # Add a column with NaN values if cleaned_text is missing

# Create and display final summary
# Ensure necessary analysis and metrics results are available before creating summary
if 'glove_analysis' in globals() and 'bert_analysis' in globals() and 'glove_metrics' in globals() and 'bert_metrics' in globals() and 'sample_df' in globals():
    results_summary = create_results_summary()
    print("\n" + "=" * 90)
    print("COMPLETE EXPERIMENT SUMMARY")
    print("-" * 90)
    print(f"\nDataset Information:")
    print(f"• Total Articles: {results_summary['dataset_info']['total_articles']}")
    print(f"• Sample Size: {results_summary['dataset_info']['sample_size']}")
    print(f"• Categories: {', '.join(results_summary['dataset_info']['categories'])}")
    print(f"• Average Text Length: {results_summary['dataset_info']['avg_text_length']:.0f} characters")
    print(f"• Average Word Count: {results_summary['dataset_info']['avg_word_count']:.0f} words")
    print(f"\nEmbedding Performance Summary:")
    # Access metrics safely in case they were not calculated
    glove_sep = results_summary['glove_results'].get('separation_score', np.nan)
    glove_coherence = results_summary['glove_results'].get('coherence_score', np.nan)
```