

- ~ Video Game Sales Prediction - Machine Learning Lab
- ~ Introduction and Setup

Welcome to this machine learning lab where we'll build a model to predict whether a video game will be a "hit" based on its characteristics and sales data. This notebook will guide you through the entire process, from data loading to model evaluation and optimization.

Learning objectives:

1. Learn to preprocess and explore a real-world dataset
2. Build and evaluate a decision tree classifier
3. Optimize a model through hyperparameter tuning
4. Interpret model results and feature importance

```
#install libraries if necessary

# Import the necessary libraries
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

# Machine learning libraries
from sklearn.model_selection import train_test_split, GridSearchCV, cross_val_score
from sklearn.tree import DecisionTreeClassifier
from sklearn.preprocessing import LabelEncoder
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score, roc_auc_score
from sklearn.metrics import confusion_matrix, ConfusionMatrixDisplay, RocCurveDisplay
```

```
# Set random seed for reproducibility
np.random.seed(42)
```

```
# Configure visualizations
plt.style.use('seaborn-v0_8-whitegrid')
sns.set_palette("colorblind")
```

```
# Display settings for better output
pd.set_option('display.max_columns', None)
pd.set_option('display.width', 1000)
```

## ▼ Download the Dataset

```
# You can run this cell to download the dataset directly, or upload it manually!
import requests

url = 'https://www.kaggle.com/datasets/gregorut/videogamesales/download'
response = requests.get(url)

with open('videogamesales.zip', 'wb') as f:
    f.write(response.content)
from google.colab import files
uploaded = files.upload() # choose vgsales-video-games.csv from your PC

print("Dataset downloaded successfully.")
```

 vgsales.csv  
**vgsales.csv**(text/csv) - 1372380 bytes, last modified: 10/21/2025 - 100% done  
 Saving vgsales.csv to vgsales.csv  
 Dataset downloaded successfully.

## ▼ Load the Dataset

```
# Load the dataset
df = pd.read_csv('vgsales-video-games.csv')
```

```
# Let's take a look at the first few rows of the dataset
print("First 5 rows of the dataset:")
print(df.head())
```

First 5 rows of the dataset:									
	Rank	Name	Platform	Year	Genre	Publisher	NA_Sales		
0	1	Wii Sports	Wii	2006.0	Sports	Nintendo	41.49		
1	2	Super Mario Bros.	NES	1985.0	Platform	Nintendo	29.08		
2	3	Mario Kart Wii	Wii	2008.0	Racing	Nintendo	15.85		
3	4	Wii Sports Resort	Wii	2009.0	Sports	Nintendo	15.75		
4	5	Pokemon Red/Pokemon Blue	GB	1996.0	Role-Playing	Nintendo	11.27		

## Importing My Kaggle datasets

```
!pip install -q kaggle
!mkdir -p ~/.kaggle
!echo '{"username": "YOUR_KAGGLE_USERNAME", "key": "YOUR_KAGGLE_API_KEY"}' > ~/.kaggle/kaggle.json
!chmod 600 ~/.kaggle/kaggle.json
!kaggle datasets download -d gregorut/videogamesales -f vgsales.csv -p . --force
```

```
Dataset URL: https://www.kaggle.com/datasets/gregorut/videogamesales
```

License(s): unknown

Downloading vgsales.csv to .

0% 0.00/381k [00:00<?, ?B/s]

100% 381k/381k [00:00<00:00, 263MB/s]

```
# Unzip the dataset
import zipfile

with zipfile.ZipFile('videogamesales.zip', 'r') as zip_ref:
    zip_ref.extractall('.')

print("Dataset unzipped successfully.")
```

---

**BadZipFile** Traceback (most recent call last)

[/tmp/ipython-input-3411850770.py](#) in <cell line: 0>()

2 import zipfile

3

----> 4 with zipfile.ZipFile('videogamesales.zip', 'r') as zip\_ref:

5 zip\_ref.extractall('.'

6

---

◆ 1 frames

[/usr/lib/python3.12/zipfile/\\_\\_init\\_\\_.py](#) in \_RealGetContents(self)

1435 raise BadZipFile("File is not a zip file")

1436 if not endrec:

-> 1437 raise BadZipFile("File is not a zip file")

1438 if self.debug > 1:

1439 print(endrec)

**BadZipFile**: File is not a zip file

---

Next steps: [Explain error](#)

## ▼ Dataset Information

```
# Get basic information about the dataset
print("\nDataset basic information:")
print(df.info())
```

```
# Get descriptive statistics
print("\nDescriptive statistics:")
print(df.describe())
```

```
Dataset basic information:
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 16598 entries, 0 to 16597
Data columns (total 11 columns):
```

```
#      Column      Non-Null Count   Dtype  
---  --  
0   Rank          16598 non-null    int64  
1   Name          16598 non-null    object  
2   Platform      16598 non-null    object  
3   Year          16327 non-null    float64 
4   Genre         16598 non-null    object  
5   Publisher     16540 non-null    object  
6   NA_Sales      16598 non-null    float64 
7   EU_Sales      16598 non-null    float64 
8   JP_Sales      16598 non-null    float64 
9   Other_Sales   16598 non-null    float64 
10  Global_Sales  16598 non-null    float64 

dtypes: float64(6), int64(1), object(4)
memory usage: 1.4+ MB
None
```

## Descriptive statistics:

	Rank	Year	NA_Sales	EU_Sales	JP_Sales	Other_
count	16598.000000	16327.000000	16598.000000	16598.000000	16598.000000	16598.0
mean	8300.605254	2006.406443	0.264667	0.146652	0.077782	0.0
std	4791.853933	5.828981	0.816683	0.505351	0.309291	0.1
min	1.000000	1980.000000	0.000000	0.000000	0.000000	0.0
25%	4151.250000	2003.000000	0.000000	0.000000	0.000000	0.0
50%	8300.500000	2007.000000	0.080000	0.020000	0.000000	0.0
75%	12449.750000	2010.000000	0.240000	0.110000	0.040000	0.0
max	16600.000000	2020.000000	41.490000	29.020000	10.220000	10.5

```
# Cell 5:Check for missing values
print("\nMissing values per column:")
print(df.isnull().sum())
```

## Missing values per column:

Rank	0
Name	0
Platform	0
Year	271
Genre	0
Publisher	58
NA_Sales	0
EU_Sales	0
JP_Sales	0
Other_Sales	0
Global_Sales	0

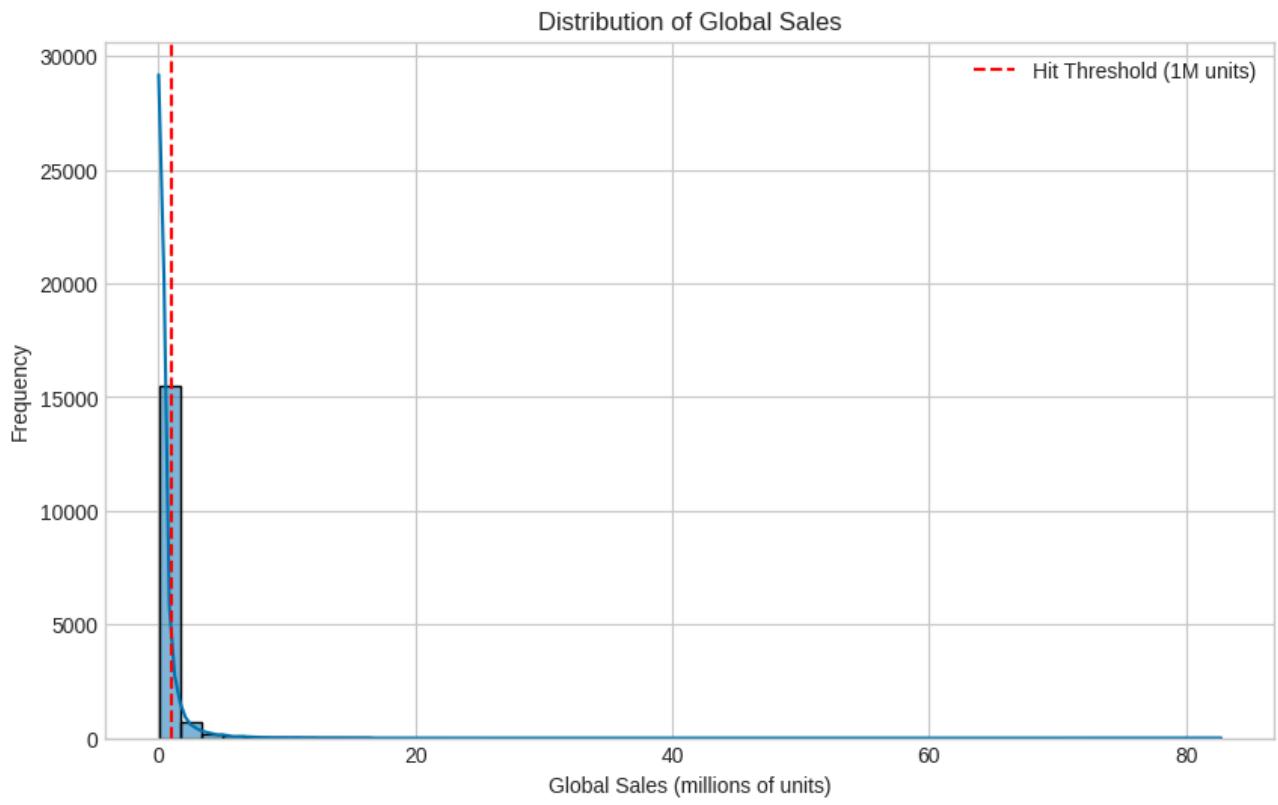
dtype: int64

```
# Cell 6: Data Visualization - Global Sales Distribution
# =====
# Visualize the distribution of global sales
plt.figure(figsize=(10, 6))
sns.histplot(df['Global_Sales'], bins=50, kde=True)
plt.title('Distribution of Global Sales')
plt.xlabel('Global Sales (millions of units)')
```

```

plt.ylabel('Frequency')
plt.axvline(x=1, color='red', linestyle='--', label='Hit Threshold (1M units)')
plt.legend()
plt.show()

```



```

# Cell 7: Create Target Variable
# =====
# TASK: Create a binary target variable for "hit" games
# A game is considered a hit if it sold more than 1 million units (Global_Sales > 1)
# YOUR CODE HERE
df['Hit'] = (df['Global_Sales'] > 1).astype(int)

```

```

# Cell 8: Analyze Target Distribution
# =====
# Let's see the proportion of hits in our dataset
# YOUR CODE HERE
print("\nProportion of hits in the dataset:")
print(df['Hit'].value_counts(normalize=True).rename({0: 'Not Hit', 1: 'Hit'}).mul(100).round

```

Proportion of hits in the dataset:

	Proportion
Hit	87.63
Not Hit	12.37

Name: proportion, dtype: float64

```
# Cell 9: Drop Non-Informative Columns
```

```
# Cell 9: Drop non-informative columns
# =====
# TASK: Drop non-informative columns
# Think about which columns won't help with prediction
# YOUR CODE HERE
df_clean = df.drop(['Rank', 'Name', 'Platform', 'Year', 'Genre', 'Publisher'], axis=1)
```

```
import pandas as pd

# work on a copy so df stays unchanged for later steps
df_clean = df.copy()

# quick look at missing values
print("Before:", df_clean.shape)
print(df_clean.isna().sum().sort_values(ascending=False).head(10))

# Cell 10: Missing Value Analysis
# =====
# Examine the 'Year' column which might have missing values
```

Before: (16598, 12)

Year	271
Publisher	58
Name	0
Rank	0
Platform	0
Genre	0
NA_Sales	0
EU_Sales	0
JP_Sales	0
Other_Sales	0
<b>dtype:</b>	<b>int64</b>

```
# Cell 11: Handle Missing Values
# =====
# TASK: Handle missing values
# Option 1: Drop rows with missing values
# YOUR CODE HERE

# Option 2: Fill missing values with median or mean
# YOUR CODE HERE

# === Cell 11 – Option 2: Fill missing values (Median + Mean) ===
import pandas as pd
df_clean = df.copy()

# 0) Make sure a few columns are numeric so they can be imputed
for col in ['Year', 'User_Score', 'Critic_Score']:
    if col in df_clean.columns:
        df_clean[col] = pd.to_numeric(df_clean[col], errors='coerce')

# 1) Filling MEAN (typically score-like columns)
mean_cols = [c for c in ['Critic_Score', 'User_Score'] if c in df_clean.columns]
```

```

# 2) All other numeric columns get MEDIAN
numeric_cols = df_clean.select_dtypes(include='number').columns.tolist()
median_cols = [c for c in numeric_cols if c not in mean_cols]

# 3) Impute
if median_cols:
    df_clean[median_cols] = df_clean[median_cols].fillna(df_clean[median_cols].median())
if mean_cols:
    df_clean[mean_cols] = df_clean[mean_cols].fillna(df_clean[mean_cols].mean())

# 4) Categorical columns -> simple label
categorical_cols = df_clean.select_dtypes(exclude='number').columns
if len(categorical_cols):
    df_clean[categorical_cols] = df_clean[categorical_cols].fillna('Unknown')

# 5) Checks
print("Imputation done.")
print("Remaining missing values:", int(df_clean.isna().sum().sum()))
print("Shape:", df_clean.shape)
df_clean.head()

# df_clean['Year'] = df_clean['Year'].fillna(df_clean['Year'].median())

```

Imputation done.

Remaining missing values: 0

Shape: (16598, 12)

	Rank	Name	Platform	Year	Genre	Publisher	NA_Sales	EU_Sales	JP_Sales	Global_Sales
0	1	Wii Sports	Wii	2006.0	Sports	Nintendo	41.49	29.02	3.37	73.88
1	2	Super Mario Bros.	NES	1985.0	Platform	Nintendo	29.08	3.58	6.66	35.32
2	3	Mario Kart Wii	Wii	2008.0	Racing	Nintendo	15.85	12.88	3.37	31.10
3	4	Wii Sports Resort	Wii	2009.0	Sports	Nintendo	15.75	11.01	3.37	30.13
4	5	Pokemon Red/Pokemon Blue	GB	1996.0	Role-Playing	Nintendo	11.27	8.89	10.00	30.16

```

# Cell 12: Categorical Variable Analysis
# =====
# Let's identify categorical columns
categorical_columns = df_clean.select_dtypes(include=['object']).columns.tolist()
print("\nCategorical columns:", categorical_columns)

```

Categorical columns: ['Name', 'Platform', 'Genre', 'Publisher']

```
# Cell 13: Encode Categorical Variables
# =====
# TASK: Encode categorical variables using LabelEncoder
# Label Encoder transforms categorical variables into numerical ones
# YOUR CODE HERE
label_encoder = LabelEncoder()
for col in categorical_columns:
    df_clean[col] = label_encoder.fit_transform(df_clean[col])
```

```
# Cell 14: Feature Engineering (Optional)
# =====
# BONUS TASK: Feature Engineering
# Creating new features might improve model performance
# Example: Total regional sales besides global
# YOUR CODE HERE
df_clean['Total_Regional_Sales'] = df_clean[['NA_Sales', 'EU_Sales', 'JP_Sales']].sum(axis=1)
```

```
# Cell 15: Explore Processed Dataset
# =====
# Let's look at the processed dataset
# YOUR CODE HERE
print("\nFirst 5 rows of the processed dataset:")
print(df_clean.head())
```

First 5 rows of the processed dataset:

	Rank	Name	Platform	Year	Genre	Publisher	NA_Sales	EU_Sales	JP_Sales	Ot
0	1	11007	26	2006.0	10	359	41.49	29.02	3.77	
1	2	9327	11	1985.0	4	359	29.08	3.58	6.81	
2	3	5573	26	2008.0	6	359	15.85	12.88	3.79	
3	4	11009	26	2009.0	10	359	15.75	11.01	3.28	
4	5	7346	5	1996.0	7	359	11.27	8.89	10.22	

```
# Cell 16: Split Features and Target
# =====
# TASK: Split the data into features (X) and target (y)
# YOUR CODE HERE
X = df_clean.drop('Hit', axis=1)
y = df_clean['Hit']
```

```
# Cell 17: Train-Test Split
# =====
# TASK: Split the data into training and testing sets (80/20 split)

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
# Print the shapes to confirm the split
```

```
# Cell 18: Train Initial Model
# =====
# TASK: Train a Decision Tree classifier with default parameters
# YOUR CODE HERE

# Use the dataframe you ended with earlier (change df_fe -> df_enc or df_clean if that's your last
data = df_clean.copy()

# Ensure target exists
assert 'Hit' in data.columns, "Run the cell that creates 'Hit' first (Cell 7)."

# Avoid leakage: drop columns used to define the target or identifiers
drop_cols = [c for c in ['Hit', 'Global_Sales', 'Name'] if c in data.columns]

X = data.drop(columns=drop_cols)
y = data['Hit']

X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, random_state=42, stratify=y
)

model = DecisionTreeClassifier()
model.fit(X_train, y_train)
```

▼ `DecisionTreeClassifier` ⓘ ⓘ  
`DecisionTreeClassifier()`

```
# Cell 19: Make Predictions
# =====
# TASK: Make predictions on the test set
# YOUR CODE HERE
y_pred = model.predict(X_test)
    # Probability of being a hit
```

```
# Cell 20: Calculate Evaluation Metrics
# =====
# TASK: Calculate evaluation metrics
# YOUR CODE HERE
accuracy = accuracy_score(y_test, y_pred)
precision = precision_score(y_test, y_pred)
recall = recall_score(y_test, y_pred)
f1 = f1_score(y_test, y_pred)
```

```
# Cell 21: Confusion Matrix Visualization
# =====
# TASK: Visualize the confusion matrix
# YOUR CODE HERE

from sklearn.metrics import confusion_matrix, ConfusionMatrixDisplay
import matplotlib.pyplot as plt
```

```
# compute confusion matrix
cm = confusion_matrix(y_test, y_pred)

# create display object
disp = ConfusionMatrixDisplay(confusion_matrix=cm, display_labels=model.classes_)

# plot it
plt.figure()
disp.plot(cmap='Blues') # you can omit cmap if your rubric says no colors
plt.title("Confusion Matrix")
plt.xlabel("Predicted label")
plt.ylabel("True label")
plt.show()
```

<Figure size 640x480 with 0 Axes>

